
NLP basics using NLTK



Contents

1. Tokenization
2. Cleaning (Stop Word Removal)
3. Stemming
4. Lemmatization
5. POS tagging
6. NER

Tokenization

Split an input text into sentence and words. You can also then use to determine bigrams and trigrams

```
from nltk.tokenize import sent_tokenize, word_tokenize  
sentence = sent_tokenize(user_input)  
words = word_tokenizer(user_input)  
  
bigrams = list(nltk.bigrams(words))  
trigrams = list(nltk.trigrams(words))
```

Stop Words Removal

Stop words are commonly used words in a language that are often filtered out or removed in natural language processing (NLP) tasks because they are deemed to be of little value in text analysis

Here are some examples of stop words in English:

- Articles: "a", "an", "the"
- Prepositions: "in", "on", "at", "by", "with"
- Conjunctions: "and", "or", "but", "so", "for"
- Pronouns: "he", "she", "it", "they", "we"

```
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize

def remove_stopwords(text):
    # Tokenize the input text
    tokens = word_tokenize(text)

    # Load the stopwords list
    stop_words = set(stopwords.words('english'))

    # Remove stopwords from the tokenized text
    filtered_tokens = [word for word in tokens if
word.lower() not in stop_words]

    # Join the filtered tokens back into a string
    filtered_text = ' '.join(filtered_tokens)

    return filtered_text
```

Stemming



Normalize words into its base form or root form

Affection

Affects

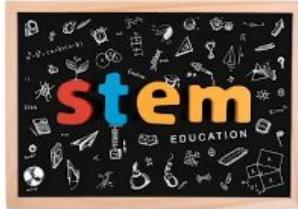
Affections

Affected

Affection

Affecting

Stemming



Normalize words into its base form or root form

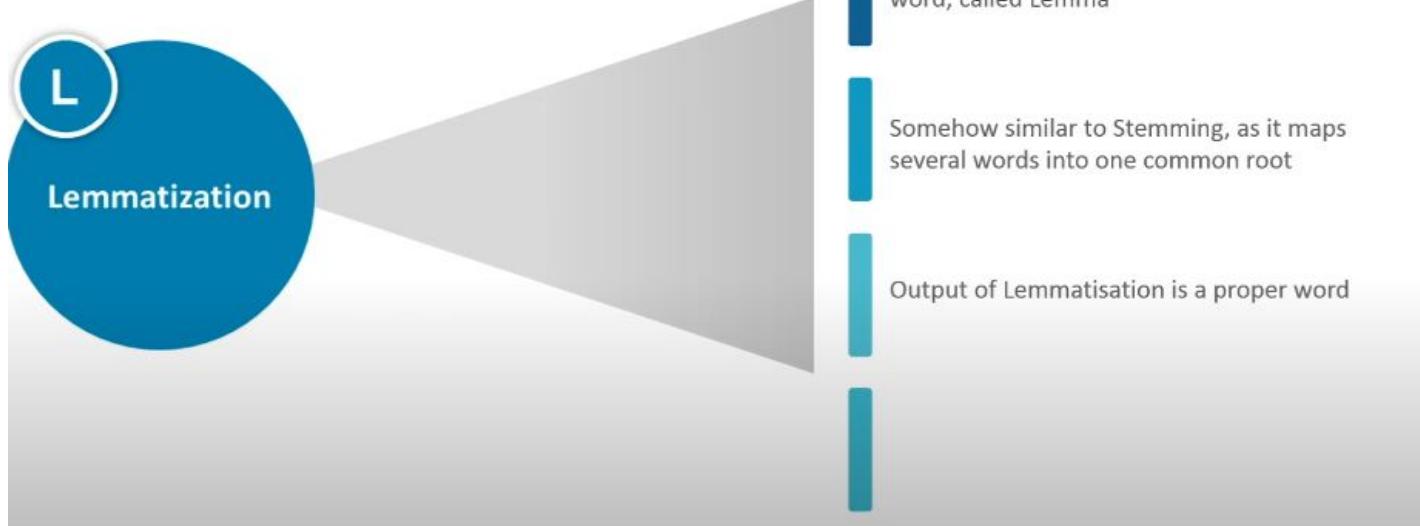
Affect

Stemming

NLTK provides various kinds of stemming algo PorterStemmer, LancasterStemmer, and SnowballStemmer

```
from nltk.stem import PorterStemmer  
pts = PorterStemmer()  
words_to_stem = ["give", "given", "gave"]  
for word in words_to_stem: print(word,':', pts.stem(word))  
  
##Output  
give: give  
giving: give  
gave:gave
```

Lemmatization





Lemmatization

```
from nltk.stem import wordnet, WordNetLemmatizer  
  
pts = WordNetLemmatizer()  
words_to_stem = ["give", "given", "gave"]  
for words in words_to_stem: print(word,':', pts.lemmatize(word))  
  
##Output  
give: give  
giving: giving  
gave:gave
```

Why no change? – By default Lemmatizer assumes POS of input word as Noun whereas all are verbs

POS: Part of Speech



POS: Tags and Description

Tag	Description
CC	Coordinating conjunction
CD	Cardinal number
DT	Determiner
EX	Existential there
FW	Foreign word
IN	Preposition or subordinating conjunction
JJ	Adjective
JJR	Adjective, comparative
JJS	Adjective, superlative
LS	List item marker
MD	Modal
NN	Noun, singular or mass
NNS	Noun, plural
NNP	Proper noun, singular
NNPS	Proper noun, plural
PDT	Predeterminer
POS	Possessive ending
PRP	Personal pronoun

Tag	Description
PRP\$	Possessive pronoun
RB	Adverb
RBR	Adverb, comparative
RBS	Adverb, superlative
RP	Particle
SYM	Symbol
TO	to
UH	Interjection
VB	Verb, base form
VBD	Verb, past tense
VBG	Verb, gerund or present participle
VBN	Verb, past participle
VBP	Verb, non3rd person singular present
VBZ	Verb, 3rd person singular present
WDT	Whdeterminer
WP	Whpronoun
WP\$	Possessive whpronoun
WRB	Whadverb

POS tagging

```
sent = "John is eating a delicious cake"
```

```
sent_tokens = word_tokenize(sent)
```

```
for token in sent_tokens: print(nltk.pos_tag(token))
```

```
##Output
```

```
[(John, NNP), [(is,VBZ), (eating, VBZ), (a, DT), (delicious, JJ), (cake, NN)]
```

NER (Named Entity Recognition)



MOVIE



MONETARY VALUE



ORGANIZATION



LOCATION



QUANTITIES



PERSON

NER

Google's CEO Sundar Pichai introduced the new Pixel at Minnesota Roi Centre Event

Organization

Person

Location

Organization



NER

```
from nltk import ne_chunk
NE_sent = "The US President stays in the White House"
NE_tokens = word_tokenize(NE_sent)
NE_tags = nltk.pos_tag(NE_tokens)
NE_NER = ne_chunk(NE_tags)

#Output
The/DT (ORG US/NNP) President/NNP stays/VBZ in/IN the/DT (Facility White/NNP House/NNP)
```

Practical Use-Case: Extractive Text Summarization

Extractive Text Summarization

Given an input document, identify top-5 sentences that captures the gist of the document. We will follow the following pipeline to generate the summary,

Segment
into
sentences

Clean
Stopwords
from
sentences

Generate
Sentence
Representations

Build Graph and
run PageRank

Generate Sentence Representation

For each sentence, we encode it to a vector using tf-idf vector representation.

- Term Frequency: $tf(t, s)$ In sentence s , the frequency of a given word t
- Document Frequency: $df(t) = \text{occurrence of } t \text{ in all sentences}$
- Inverse Document Frequency: $idf(t) = |S|/df(t)$

$$tf\text{-}idf(t, d) = tf(t, d) * idf(t)$$

However you do not have to write this from scratch, we will use scikit-learn

TF-IDF using scikit-learn

```
# import required module  
from sklearn.feature_extraction.text import TfidfVectorizer  
  
# assign sentences  
d0 = 'Geeks for geeks'  
d1 = 'Geeks'  
d2 = 'r2j'
```

Word indexes:

```
{'geeks': 1, 'for': 0, 'r2j': 2}
```

```
# merge sentences into a single corpus  
string = [d0, d1, d2]
```

```
# create object  
tfidf = TfidfVectorizer()
```

```
# get tf-df values  
result = tfidf.fit_transform(string)
```

tf-idf values in matrix form:

```
[[0.54935123 0.83559154 0.  
 [0. 1. 0.  
 [0. 0. 1.  
 ]]  
 ]]
```

Pagerank-based summarization

Create a fully-connected graph using sentences as nodes and tf-idf similarity as edge weights

Run Pagerank to identify important nodes (Sentences)

Select top-5 nodes (sentences) as the summary.