



Systems Programming: Threading (Part 1)

Computing Lab (cs 69201)
Department of Computer Science and Engineering
IIT Kharagpur

What is systems program?



- Programming that requires close interaction with the operating system

What is systems program?



- Programming that requires close interaction with the operating system
- Our target areas:
 - Operating Systems
 - Networking Software
 - Compilers & Interpreters ... so on.



What is systems program?

- Programming that requires close interaction with the operating system
- Our target areas:
 - Operating Systems
 - Networking Software
 - Compilers & Interpreters ... so on.
- Why are we even studying it ?
 - Forms the foundation of all software
 - Critical for performance-sensitive applications
 - Enables efficient use of hardware resources

Quick Recap



- **Processes**

- Independent units of execution
- Have their own memory space
- Created through system calls (e.g., `fork()` in Unix)
- Communicate via Inter-Process Communication (IPC)

Quick Recap



- **Processes**

- Independent units of execution
- Have their own memory space
- Created through system calls (e.g., fork() in Unix)
- Communicate via Inter-Process Communication (IPC)

- **Inter-Process Communication (IPC)**

- Methods for processes to exchange data
- Includes pipes, shared memory,

Introduction to Concurrency



- **What is Concurrency?**
 - Ability to handle multiple tasks simultaneously. **Sure ?**

Introduction to Concurrency



- **What is Concurrency?**
 - Ability to handle multiple tasks simultaneously. **Sure ?**
- **Need of Concurrency in System Programming?**
 - Utilizes multi-core processors effectively
 - Improves system responsiveness
 - Enables parallel processing of tasks

Introduction to Concurrency



- **What is Concurrency?**
 - Ability to handle multiple tasks simultaneously. **Sure ?**
- **Need of Concurrency in System Programming?**
 - Utilizes multi-core processors effectively
 - Improves system responsiveness
 - Enables parallel processing of tasks
- **Challenges of Concurrency**
 - Race conditions
 - Deadlocks
 - Increased complexity in program design

Is concurrency **same** as Parallelism?



Is concurrency same as Parallelism?

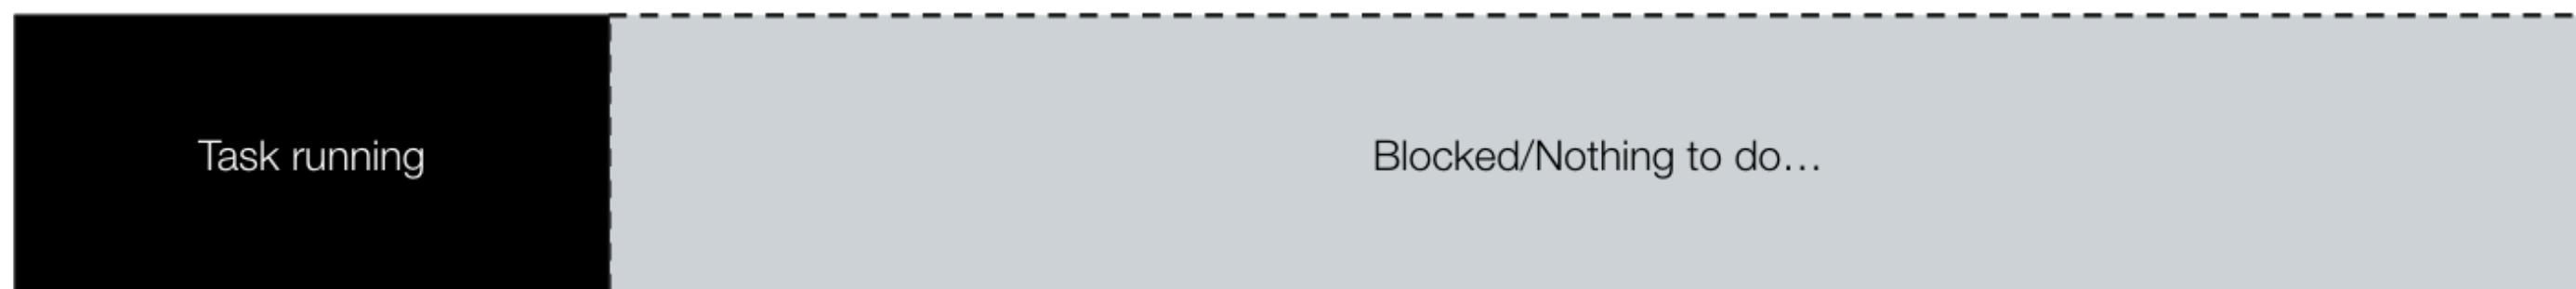


- Parallelism is when tasks literally run at the same time, e.g., on a multicore processor.
- Parallelising CPU bound tasks improves throughput.



Is concurrency same as Parallelism?

- Parallelism is when tasks literally run at the same time, e.g., on a multicore processor.
- Parallelising CPU bound tasks improves throughput.
- But what if, tasks are blocked/idle most of the time?

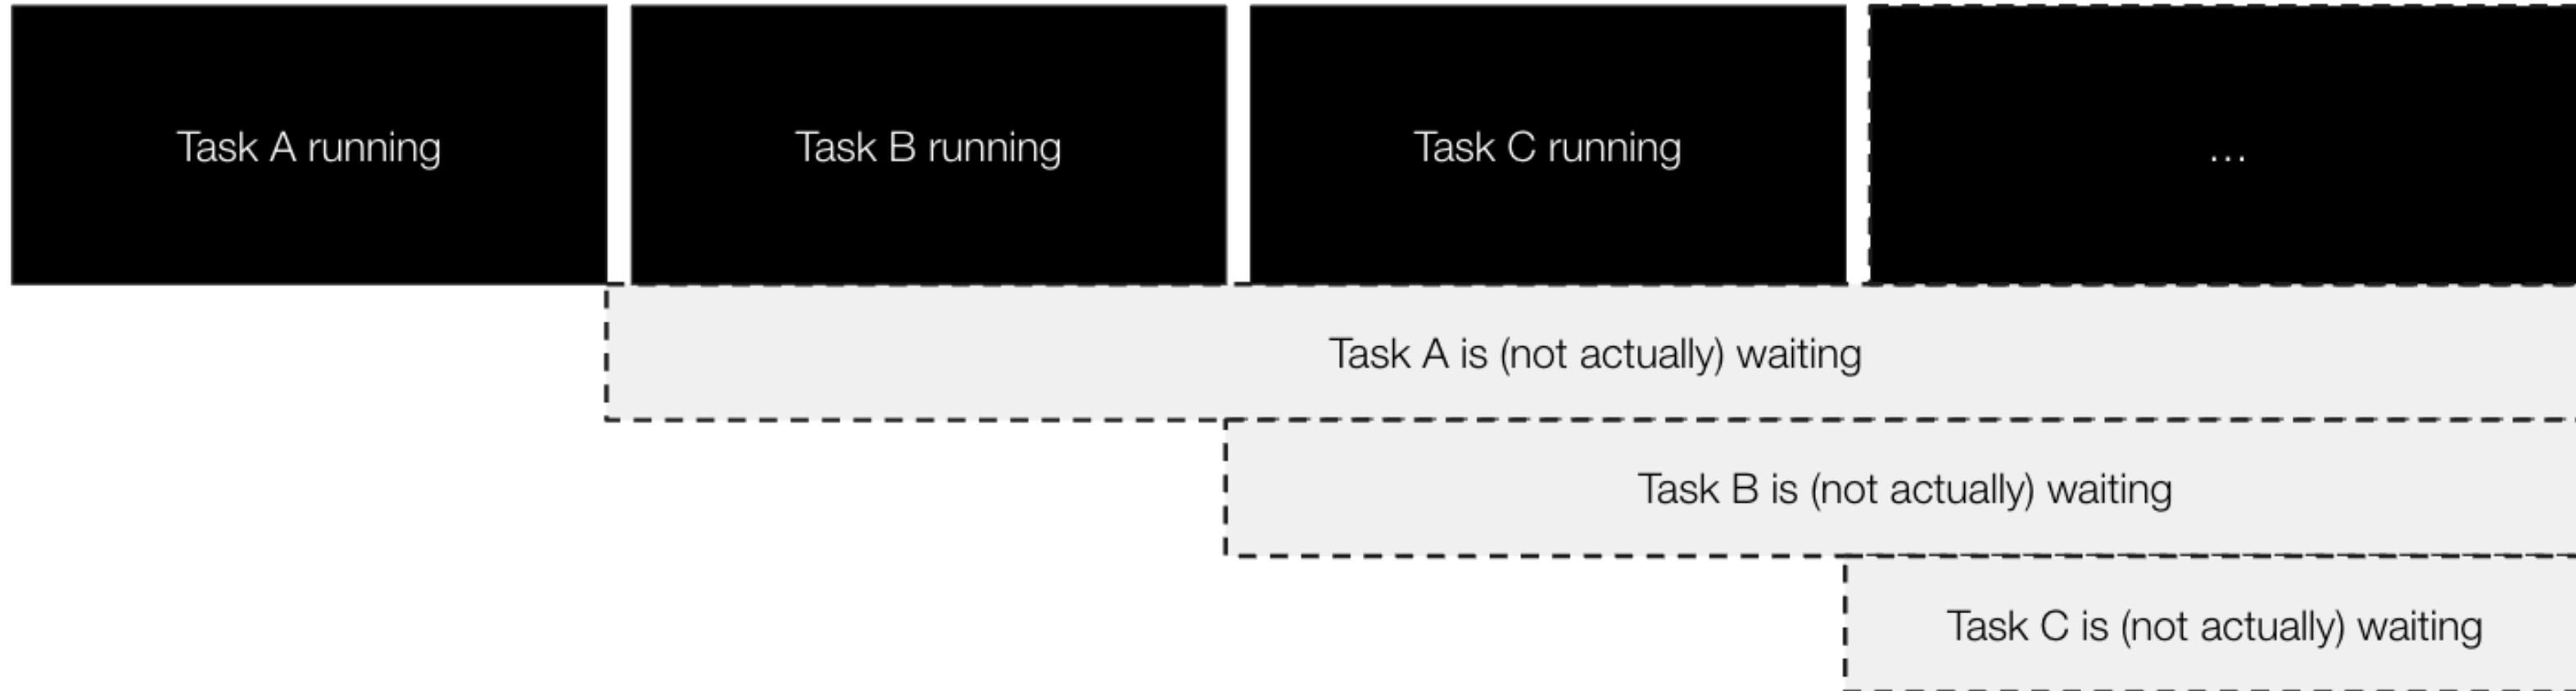


So parallelise it ?

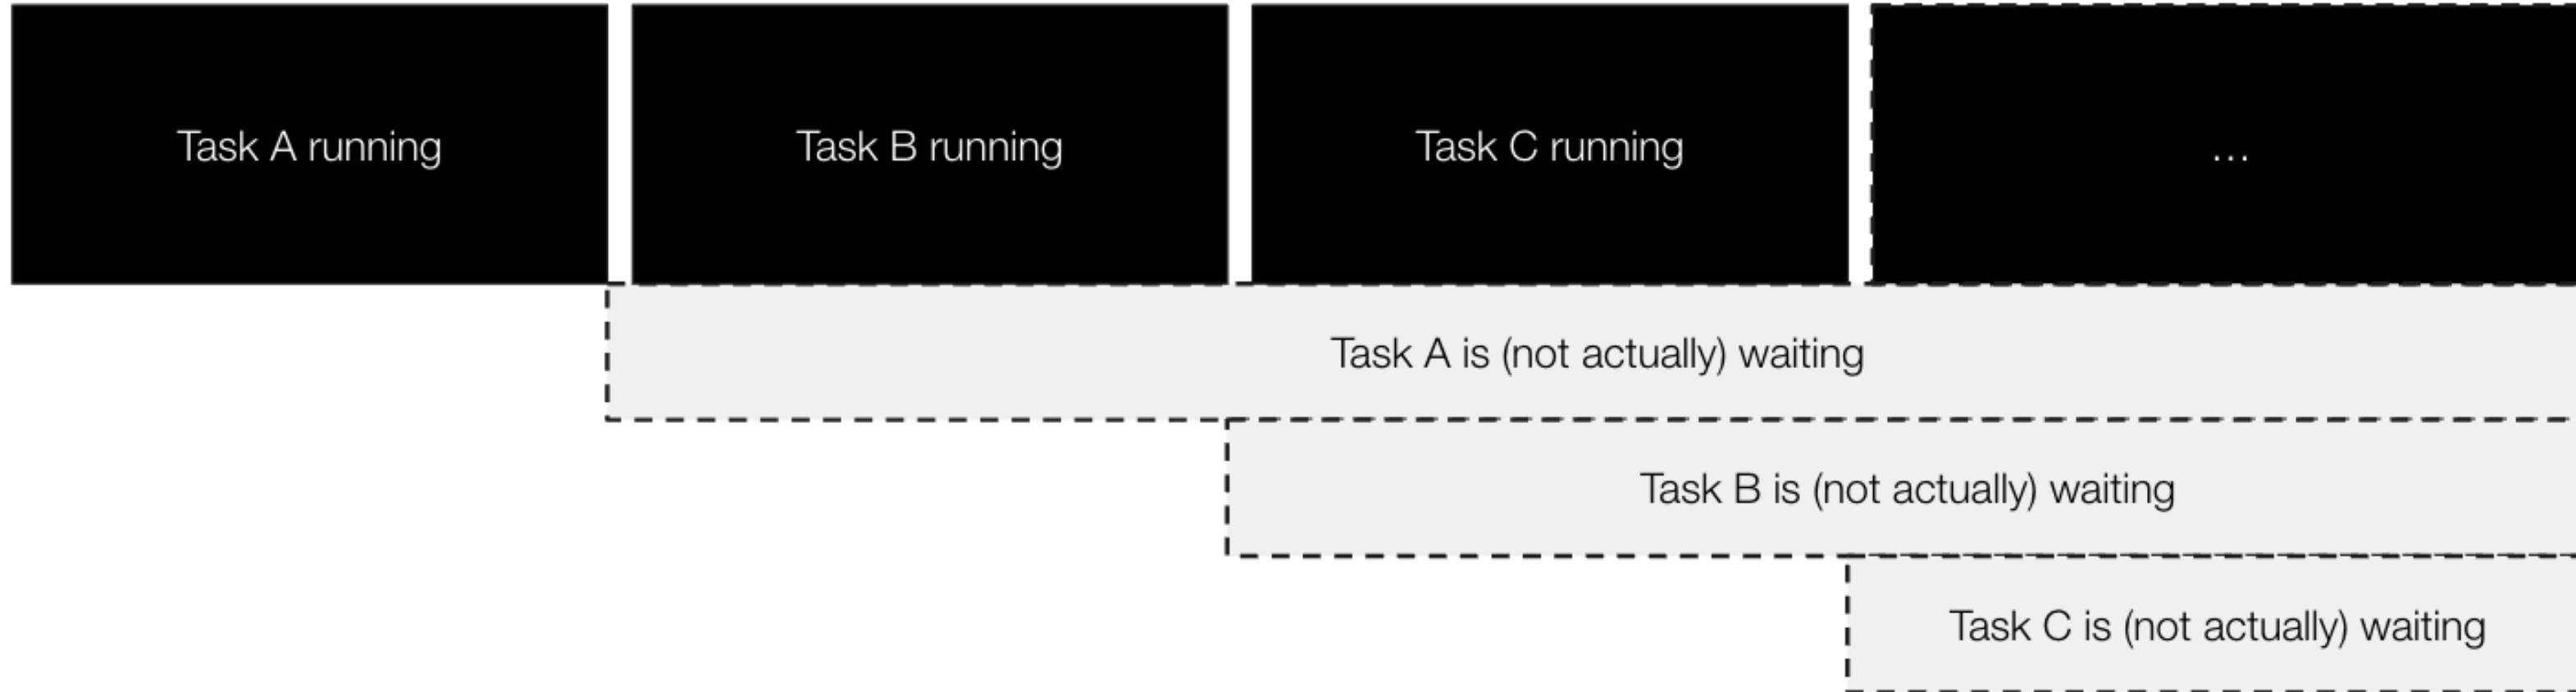


- It gets a bit better... but clearly this can't be the best.
- We can do **better**

The idea of concurrency

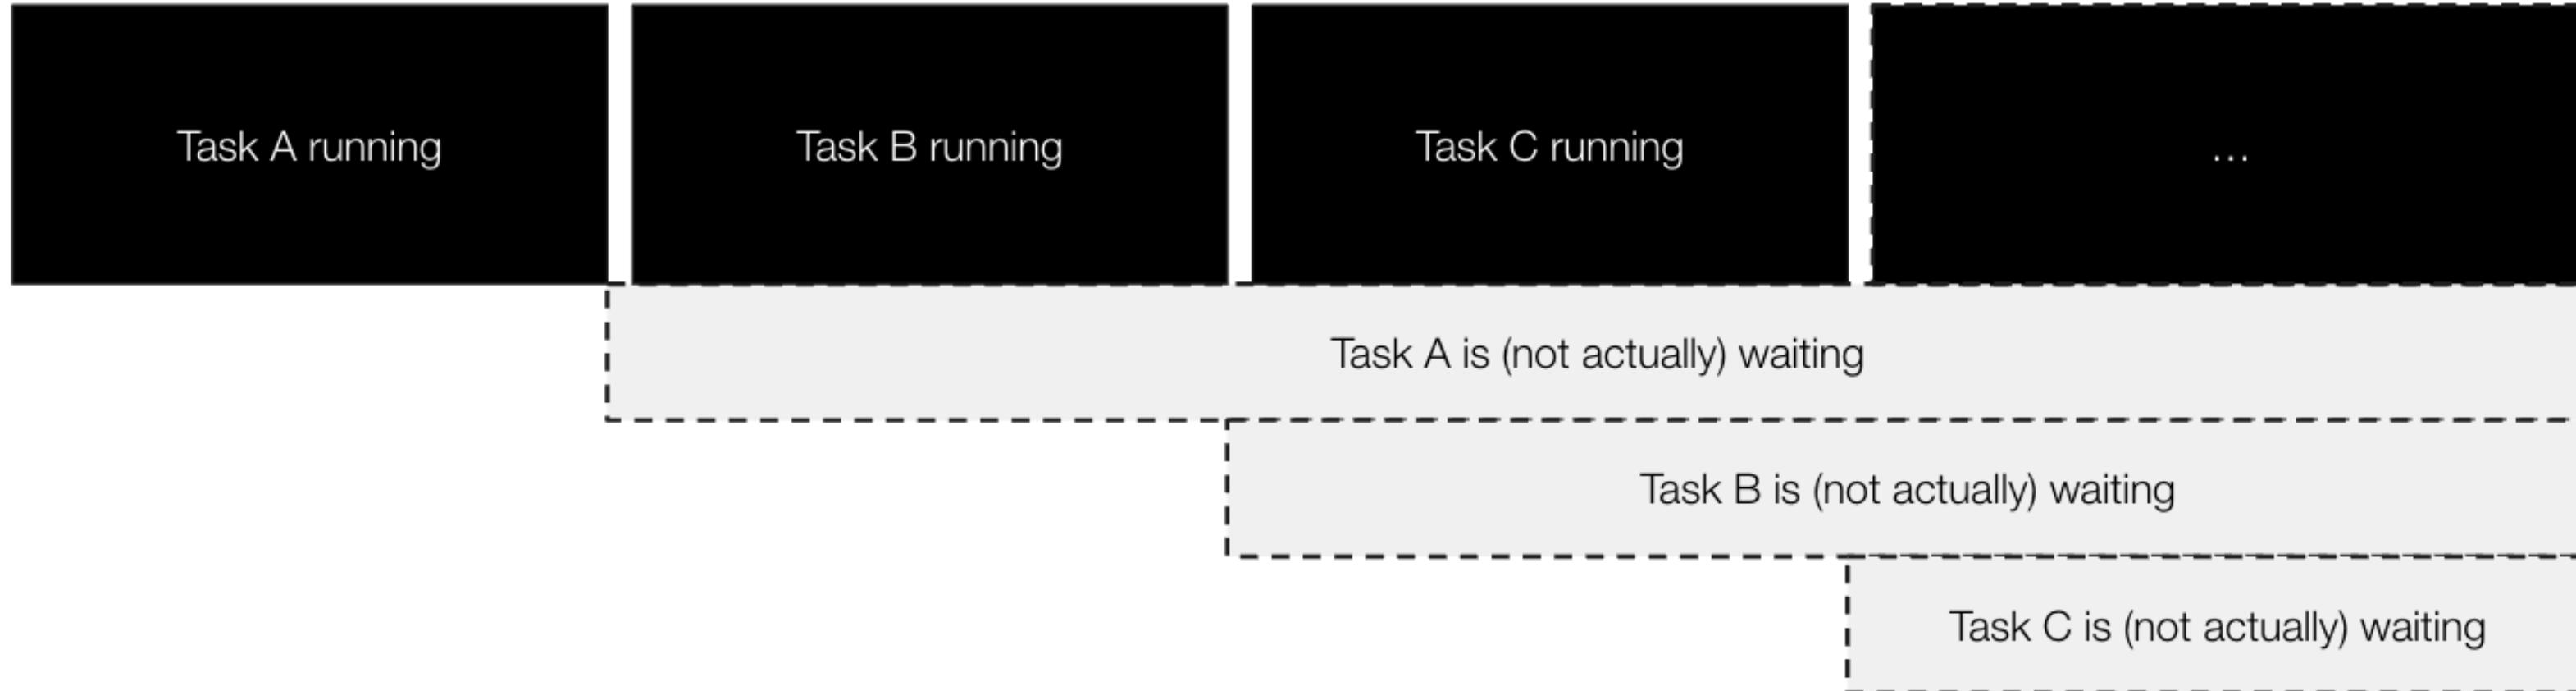


The idea of concurrency



- Concurrency is when two or more tasks can start, run, and complete together [doesn't necessarily mean they'll ever both be running at the same instant]
- For example, multitasking on a single-core machine.
- Concurrent IO bound and CPU bound tasks improve throughput

The idea of concurrency



- Concurrency is when two or more tasks can start, run, and complete together [doesn't necessarily mean they'll ever both be running at the same instant]
- For example, multitasking on a single-core machine.
- Concurrent IO bound and CPU bound tasks improve throughput

do you find any catch here ?



This is where the fun begins.



How to create a process ?



How to create a process ?



```
● ● ●

#include <stdio.h>
#include <unistd.h>

int main() {
    // Create a new process using fork()
    pid_t pid = fork();

    if (pid == 0) {
        // Child process: fork() returns 0
        printf("This is the child process with PID: %d\n", getpid());
    } else if (pid > 0) {
        // Parent process: fork() returns the PID of the child
        printf("This is the parent process with PID: %d\n", getpid());
        printf("Created a child process with PID: %d\n", pid);
    } else {
        // Fork failed
        printf("Failed to create a new process\n");
    }
}

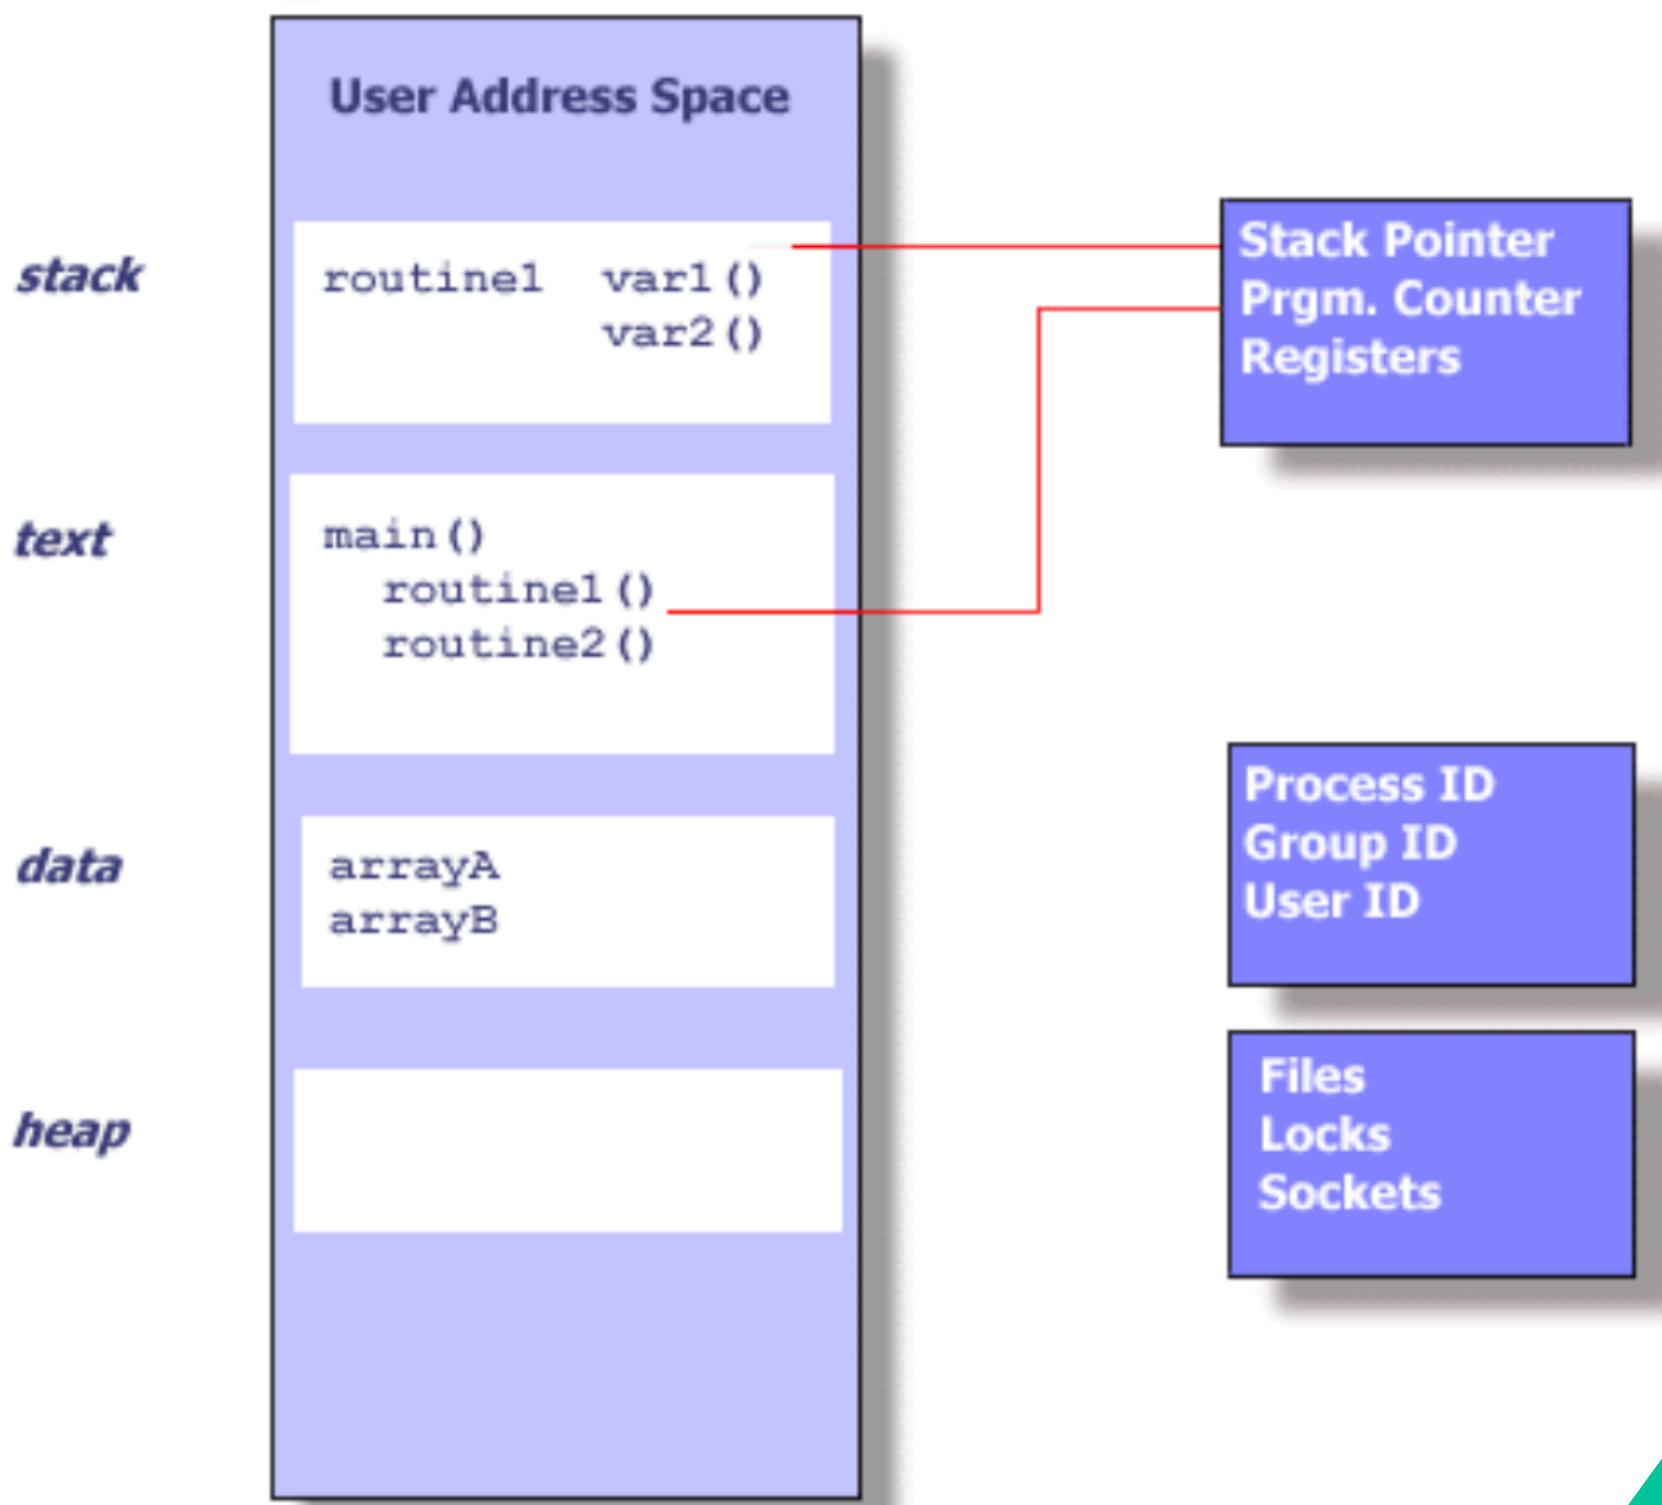
return 0;
}
```

Process & Threads (1/4)



Processes contain information about program resources and program execution state, including:

- Process ID, process group ID, user ID, and group ID, address space
- Program instructions, registers, stack, heap.
- File descriptors, inter-process communication tools
- Shared libraries



Process & Threads

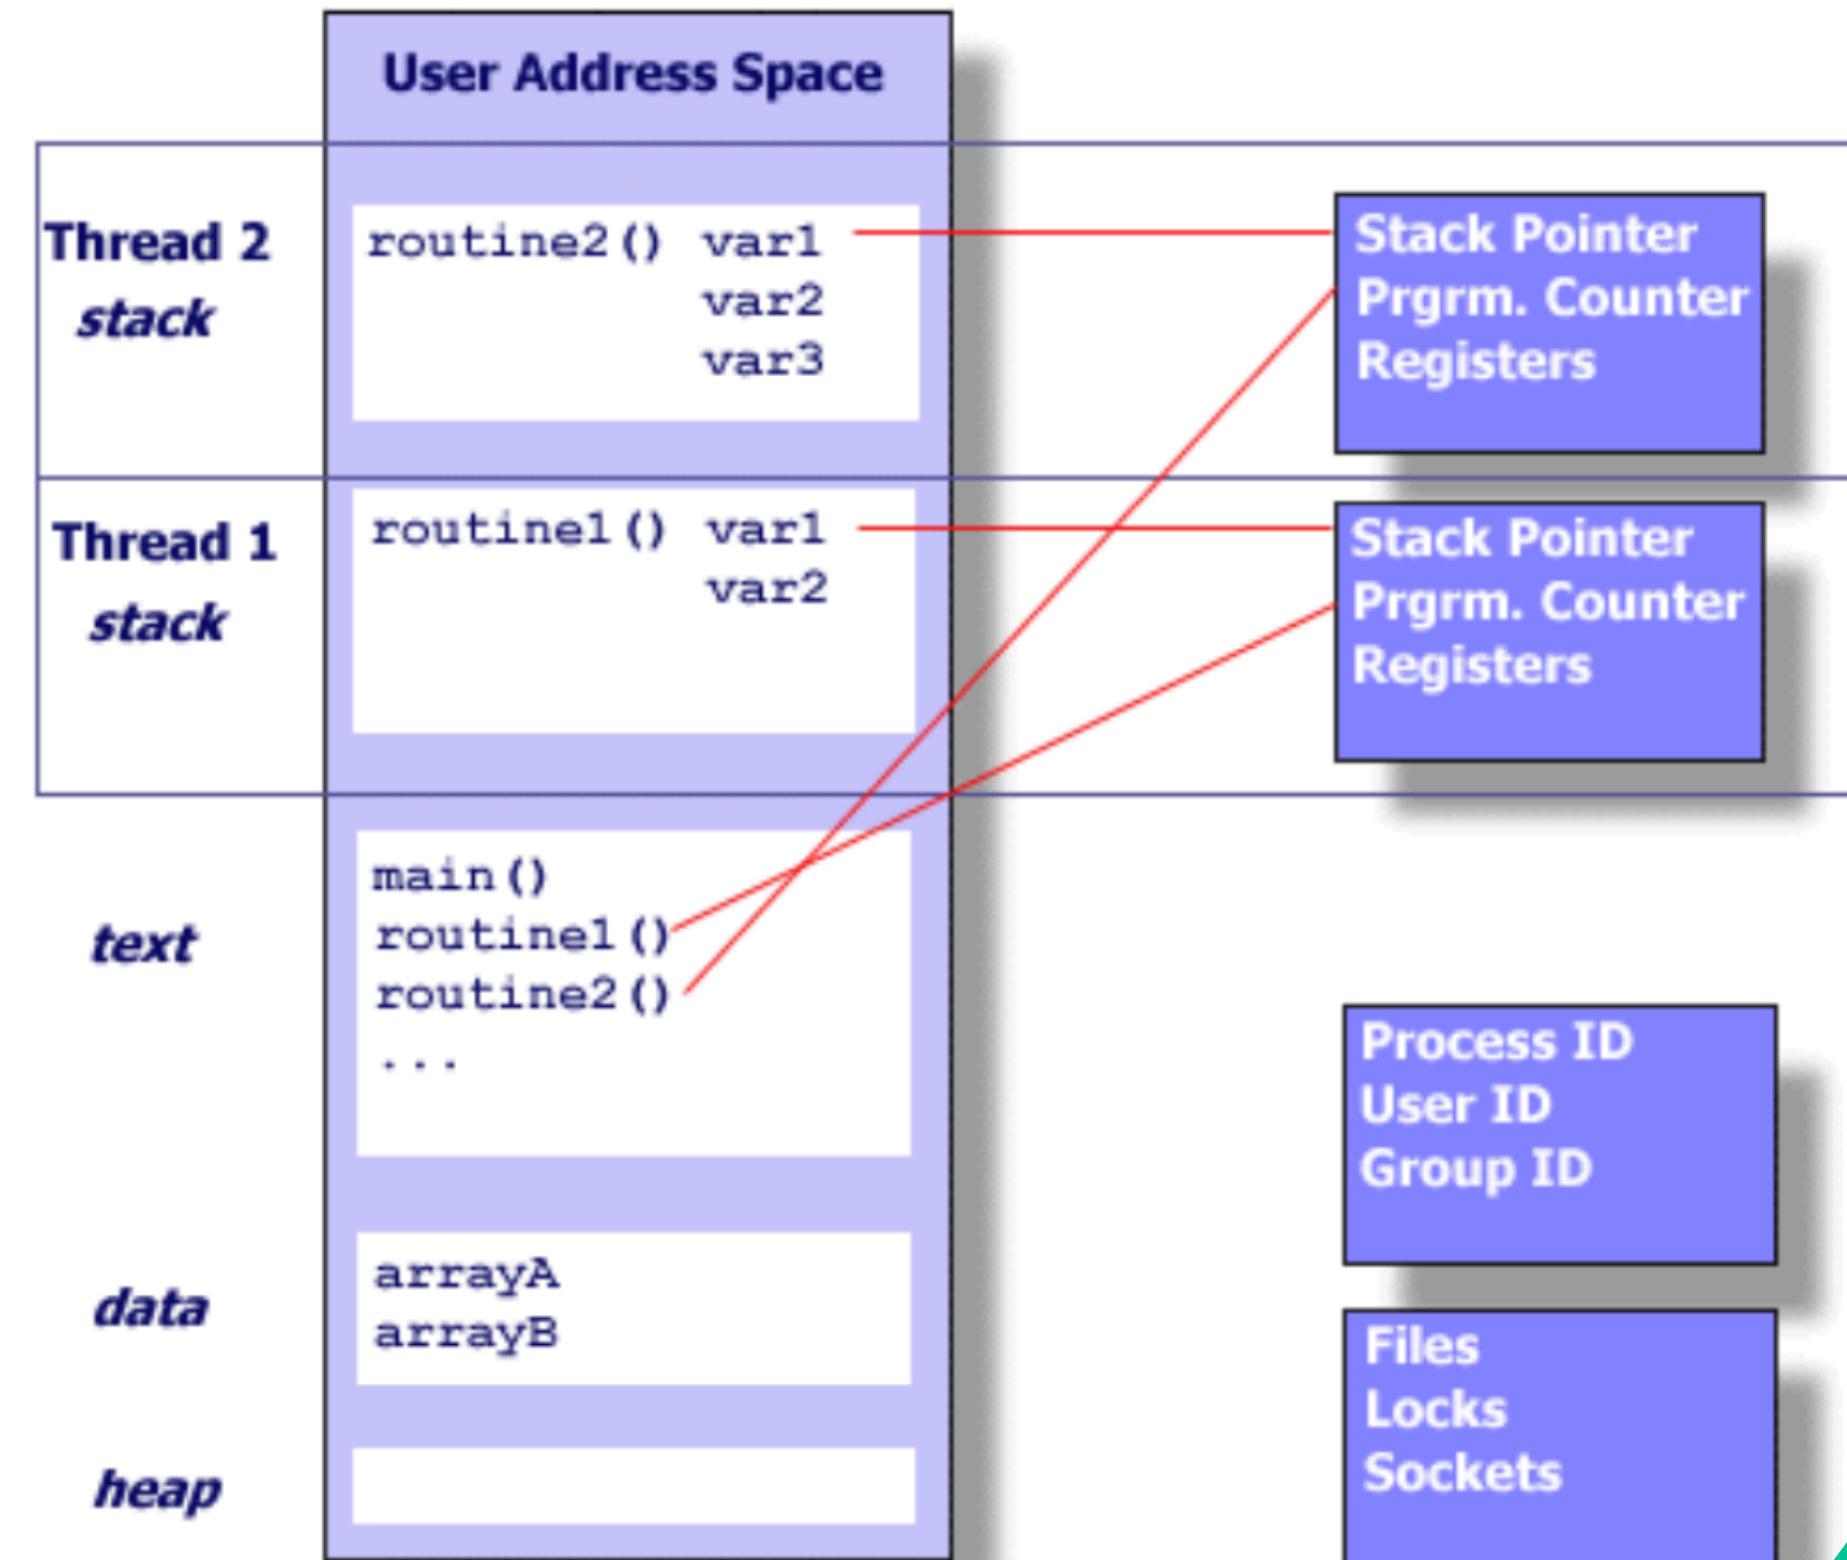


```
#include <stdio.h>
#include <unistd.h>

int main() {
    printf("Process ID: %d\n", getpid());           // Get Process ID
    printf("Parent Process ID: %d\n", getppid());     // Get Parent Process ID
    printf("User ID: %d\n", getuid());                // Get User ID
    printf("Group ID: %d\n", getgid());                // Get Group ID
    return 0;
}
```

Process & Threads (2/4)

Threads use and exist within these process resources, yet are able to be scheduled by the operating system and run as independent entities within a process



Process & Threads



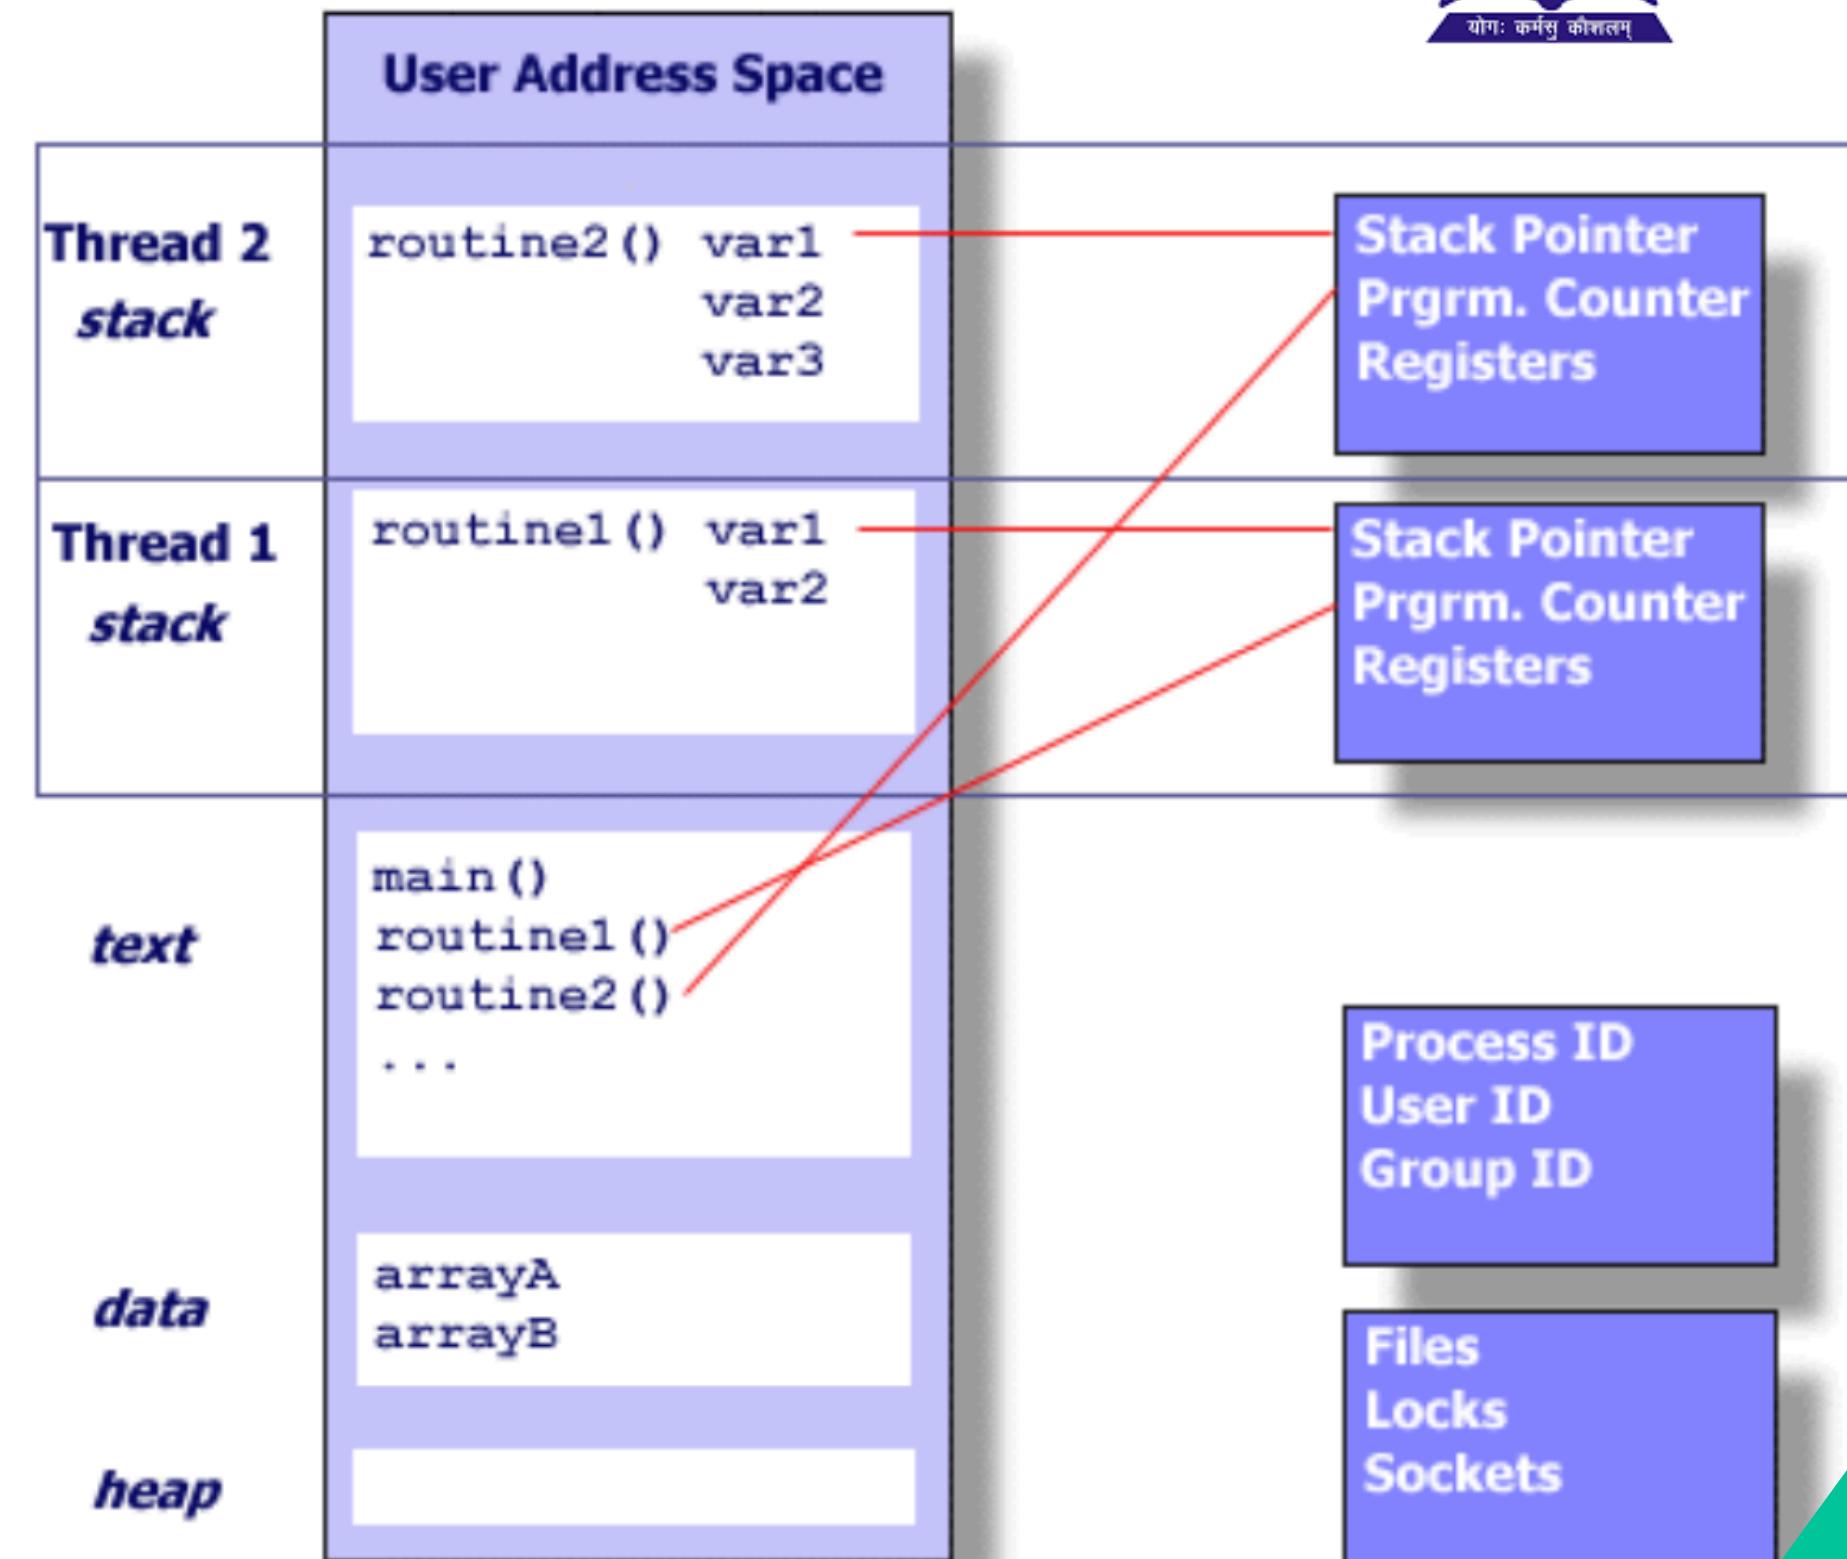
```
#include <stdio.h>
#include <pthread.h>

void* thread_function(void* arg) {
    printf("Thread running with ID: %ld\n", pthread_self());
    return NULL;
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, thread_function, NULL); // Create a new
thread
    pthread_join(thread, NULL); // Wait for thread to finish
    printf("Main process finished\n");
    return 0;
}
```

Process & Threads (3/4)

- A thread can possess an independent flow of control and be schedulable because it maintains its own:
 - Stack pointer
 - Registers
 - Scheduling properties
 - Set of pending and blocked signals
 - Thread specific data.



Process & Threads (4/4)

Because threads within the same process share resources:

- Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads .
- Two pointers having the same value point to the same data.
- Reading and writing to the same memory locations is possible, and therefore requires explicit synchronization by the programmer



Process & Threads



```
#include <stdio.h>
#include <pthread.h>

int shared_resource = 0;

void* thread_function(void* arg) {
    shared_resource++;
    printf("Thread updated shared resource to %d\n", shared_resource);
    return NULL;
}

int main() {
    pthread_t thread1, thread2;
    pthread_create(&thread1, NULL, thread_function, NULL);
    pthread_create(&thread2, NULL, thread_function, NULL);

    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);

    printf("Final value of shared resource: %d\n", shared_resource);
    return 0;
}
```

pthreads



- For UNIX systems, implementations of threads that adhere to the IEEE POSIX 1003.1c standard are Pthreads.

pthreads



- For UNIX systems, implementations of threads that adhere to the IEEE POSIX 1003.1c standard are Pthreads.
- The primary motivation behind Pthreads is improving program performance
- Can be created with much less OS overhead & Needs fewer system resources to run.

pthread library



- Programs must include the file pthread.h
- Programs must be linked with the pthread library
 - gcc -lpthread main.c -o main
- Some functions
 - pthread_create(): Creates a new thread
 - pthread_exit(): Terminates the calling thread
 - pthread_mutex_lock(): lock a mutex
 - pthread_join(): Causes the calling thread to wait for another thread to terminate

pthread library



- Programs must include the file pthread.h
- Programs must be linked with the pthread library
 - gcc -lpthread main.c -o main
- Some functions
 - pthread_create(): Creates a new thread
 - pthread_exit(): Terminates the calling thread
 - pthread_mutex_lock(): lock a mutex
 - pthread_join(): Causes the calling thread to wait for another thread to terminate

explore the functions !!!

Example: Creating a thread with Pthreads



```
#include <pthread.h>

void *thread_function(void *arg) {
    // Thread code here
    return NULL;
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, thread_function, NULL);
    pthread_join(thread, NULL);
    return 0;
}
```

C++ std::thread

- Provides a higher-level abstraction over native threads



C++ std::thread

- Provides a higher-level abstraction over native threads
- Automatic resource management
- Integrates well with other C++ features (lambdas, function objects)
- Easier to write clean code



C++ std::thread

- Provides a higher-level abstraction over native threads
- Automatic resource management
- Integrates well with other C++ features (lambdas, function objects)
- Easier to write clean code
- Disadvantages:
 - Not all synchronization primitives are implemented, e.g., barriers (do not worry, you will not implement it this week).
 - A modern compiler is needed



Thread creation



```
#include <thread>
#include <iostream>
#include <string>

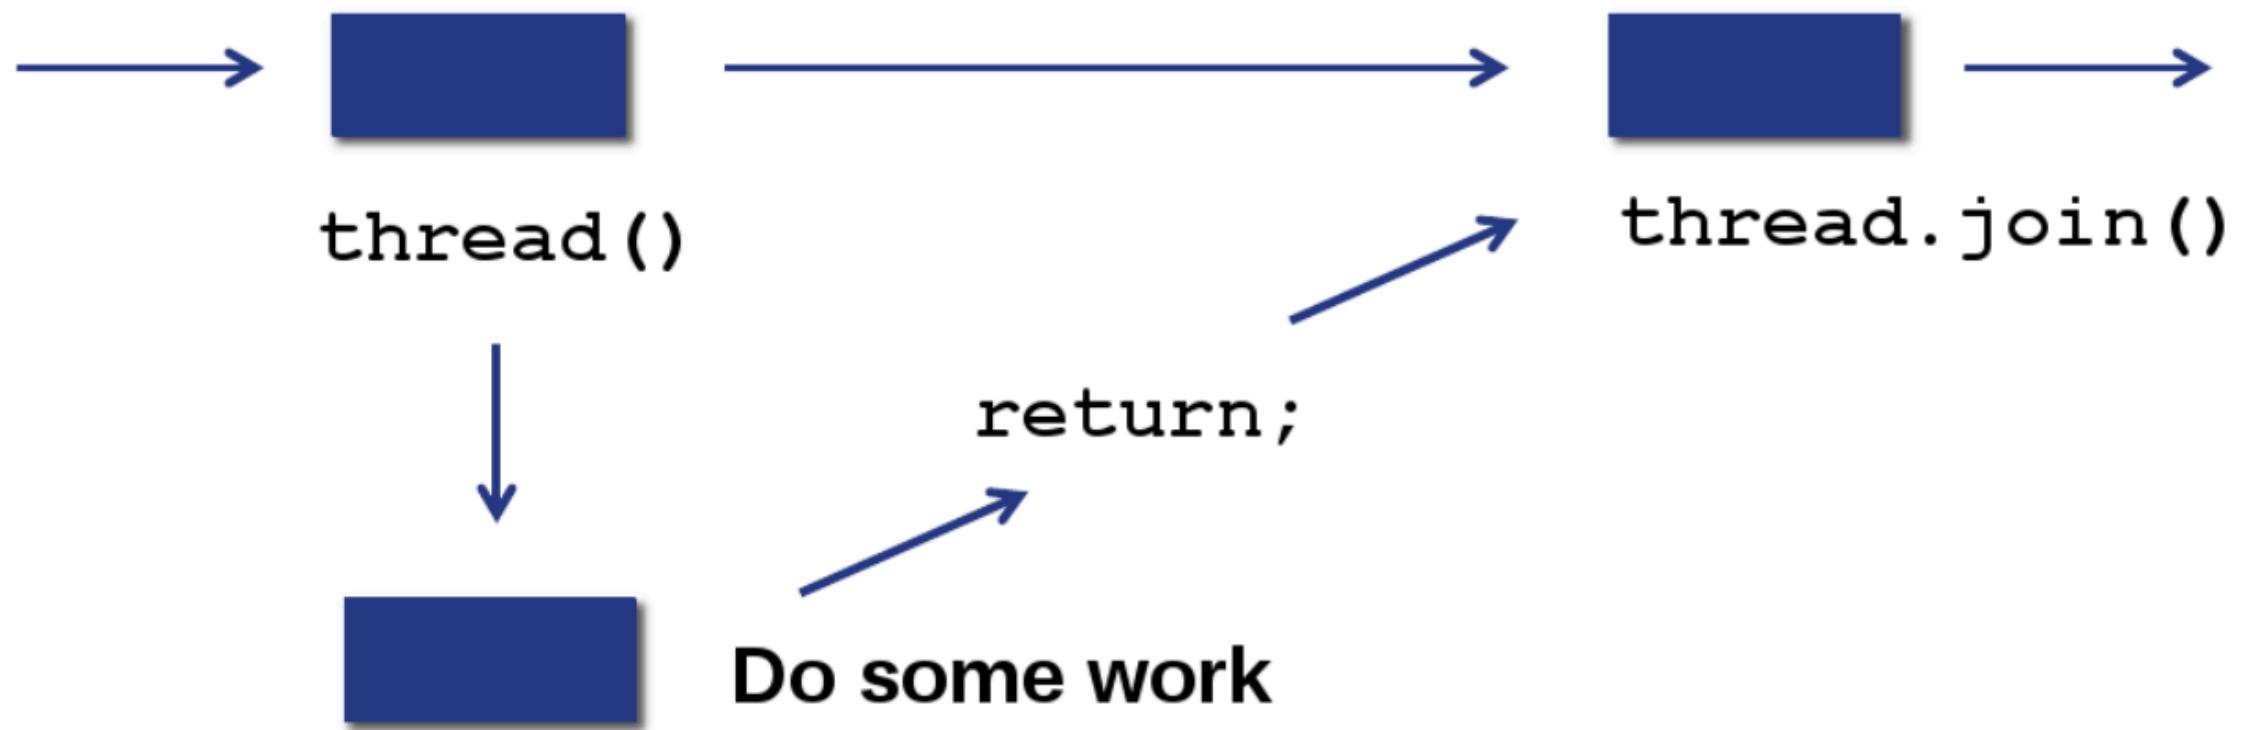
void greet(const std::string& name) {
    std::cout << "Hello, " << name << "!" << std::endl;
}

int main() {
    std::thread t(greet, "Alice");
    t.join();
    return 0;
}
```



Joining of threads

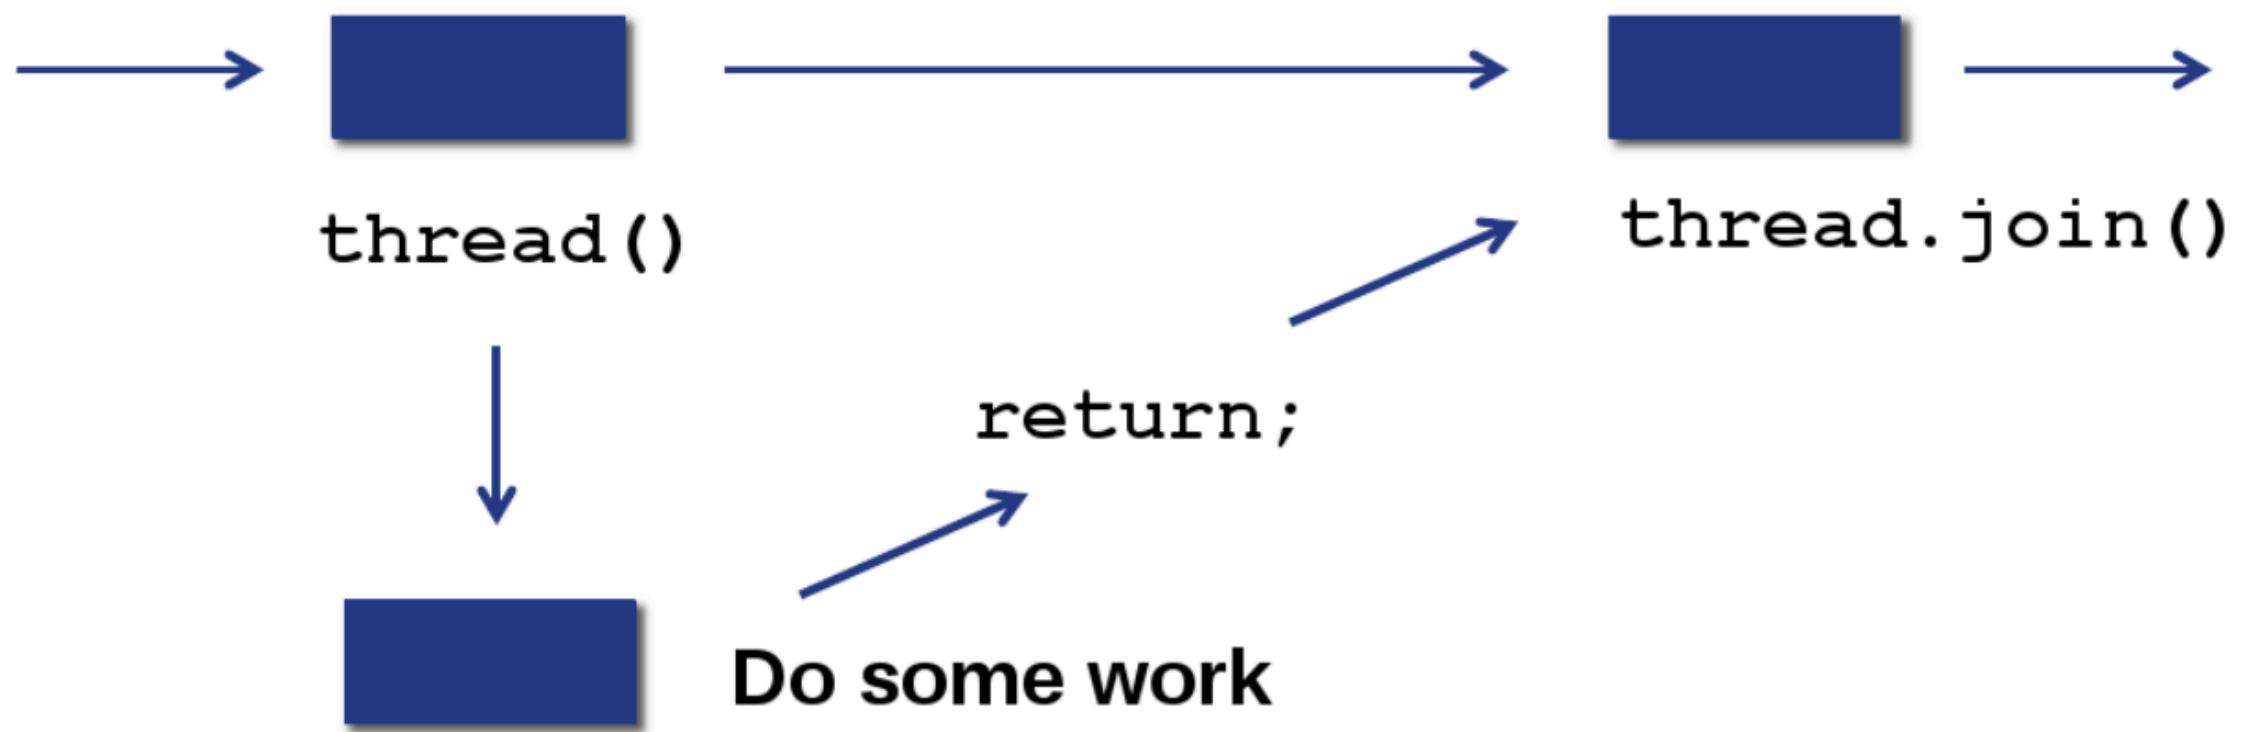
- `join()`: Wait for thread to finish





Joining of threads

- `join()`: Wait for thread to finish



```
std::thread t(some_function);
t.join(); // Wait for t to finish
```

Can you find out the problem in the above implementation ?



Can you find out the problem in the above implementation ?



- Do you all remember about the catch in OS threads ?

Latency & Overhead

Can you find out the problem in the above implementation ?



- Do you all remember about the catch in OS threads ?

Latency & Overhead

- Starting up an OS thread, or even waking it up takes many clock cycles
- Hence, OS threads should not be used when compute time is small, and is overshadowed by creation/start-up time

Whats the solution ?

- We need to create infrastructure for task loading and unloading the the threads.



Whats the solution ?

- We need to create infrastructure for task loading and unloading the the threads.
- And thats called an “**EVENT LOOP**”



Whats the solution ?

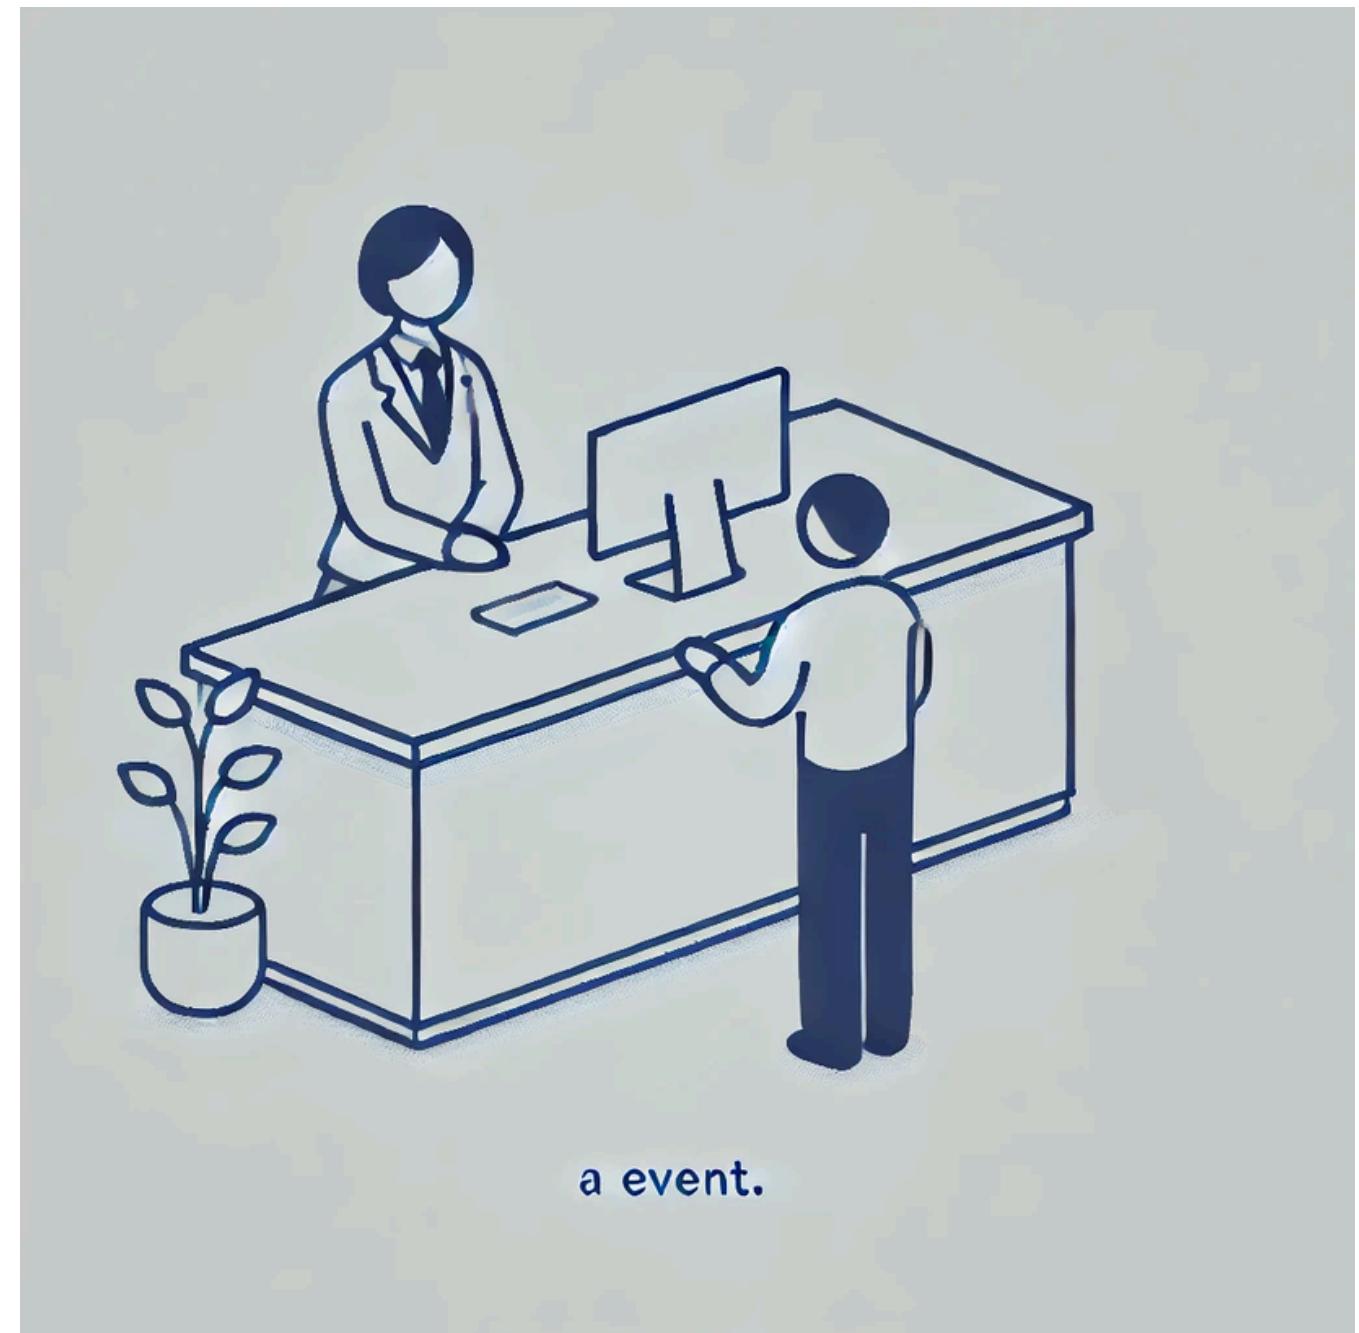


- We need to create infrastructure for task loading and unloading the threads.
- And that's called an “**EVENT LOOP**”
- And what is an event ?
 - Events can be things like keyboard input, mouse clicks, or network data (like interrupts)



Analogy to an event loop

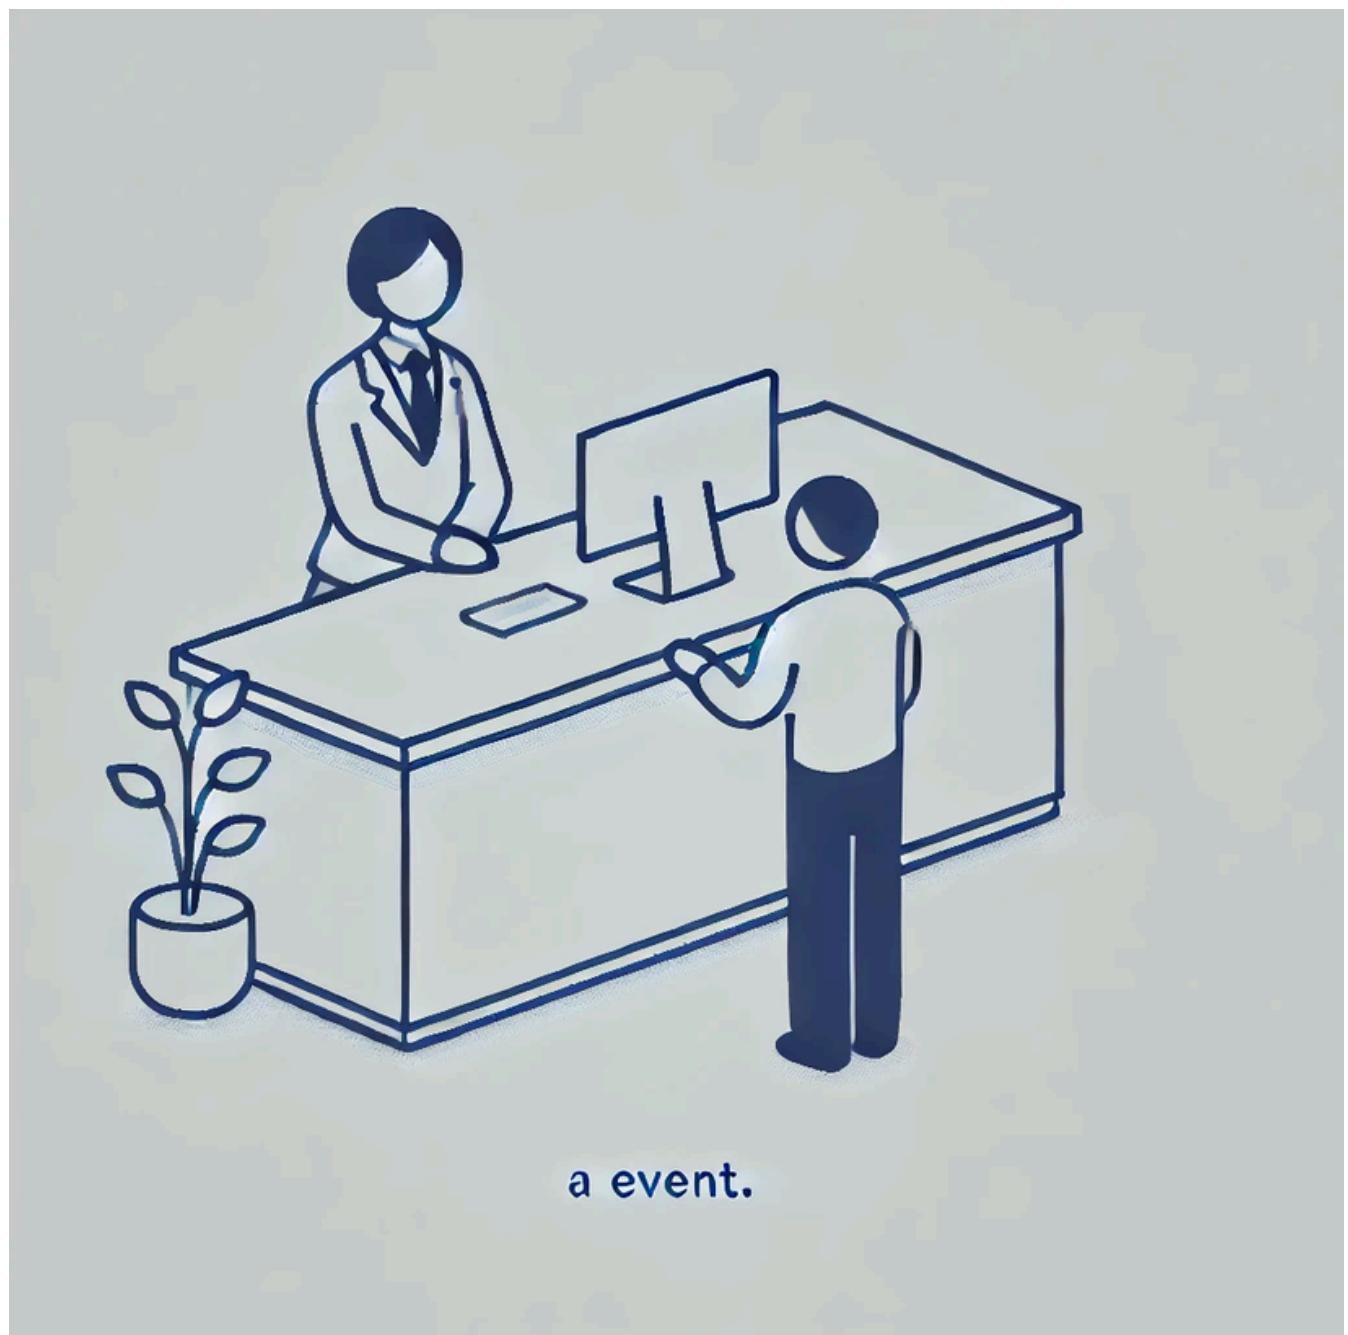
- The receptionist waits for someone (**event**) to walk in.





Analogy to an event loop

- The receptionist waits for someone (**event**) to walk in.
- When someone comes in with a question (**event**), the receptionist helps them (**processes the event**).

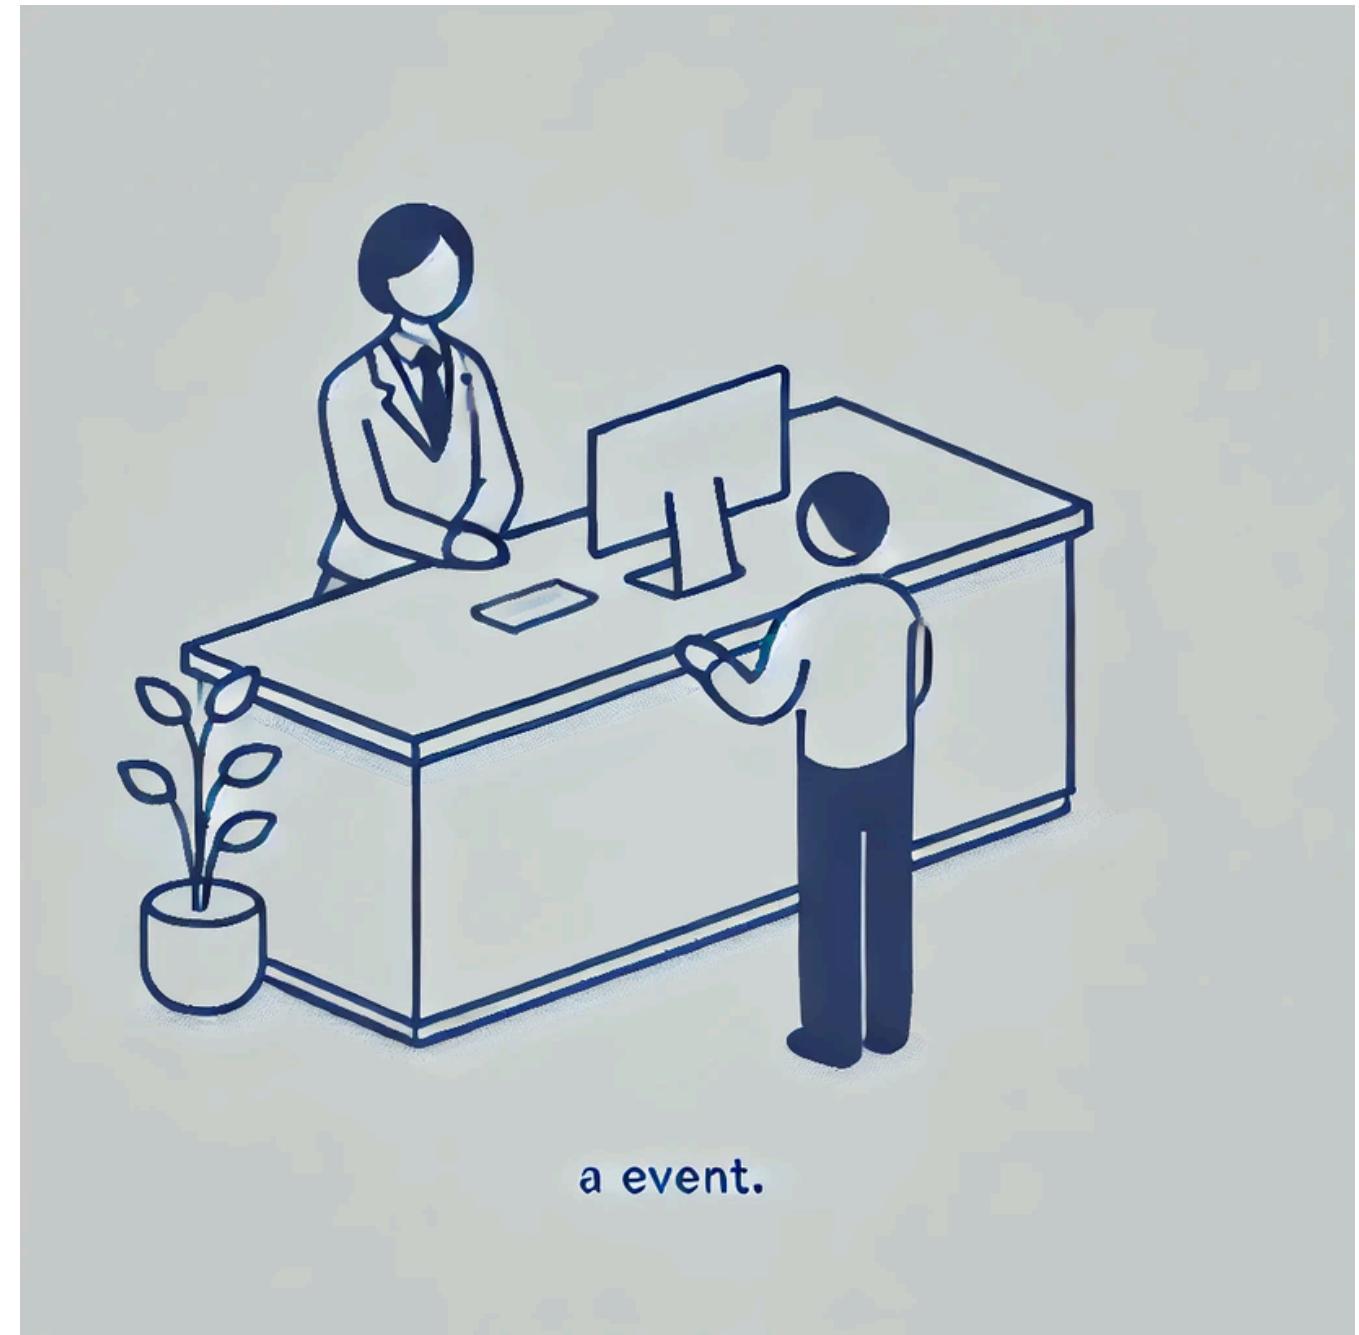


a event.



Analogy to an event loop

- The receptionist waits for someone (**event**) to walk in.
- When someone comes in with a question (**event**), the receptionist helps them (**processes the event**).
- Once that person is helped, the receptionist is ready for the next person. (**event**)



a event.

So what is an event loop ?

- An event loop in an operating system is like a manager that listens for events.



So what is an event loop ?

- An event loop in an operating system is like a manager that listens for events.
- When an event happens, the event loop processes it by triggering the appropriate action (function or handler).



So what is an event loop ?

- An event loop in an operating system is like a manager that listens for events.
- When an event happens, the event loop processes it by triggering the appropriate action (function or handler).
- After processing the event, the event loop goes back to listening for more events.



So what is an event loop ?

- An event loop in an operating system is like a manager that listens for events.
- When an event happens, the event loop processes it by triggering the appropriate action (function or handler).
- After processing the event, the event loop goes back to listening for more events.
- It allows the system to handle multiple tasks without getting stuck on one.



How to implement concurrency?

Step-1

- Start a task A, and do required computation

Task A running

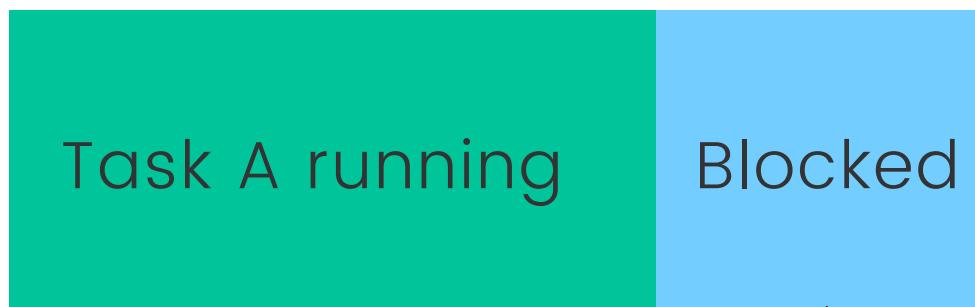


How to implement concurrency?

Step-2



- Start a task A, and do required computation
- When task A blocks/waits, save state of A



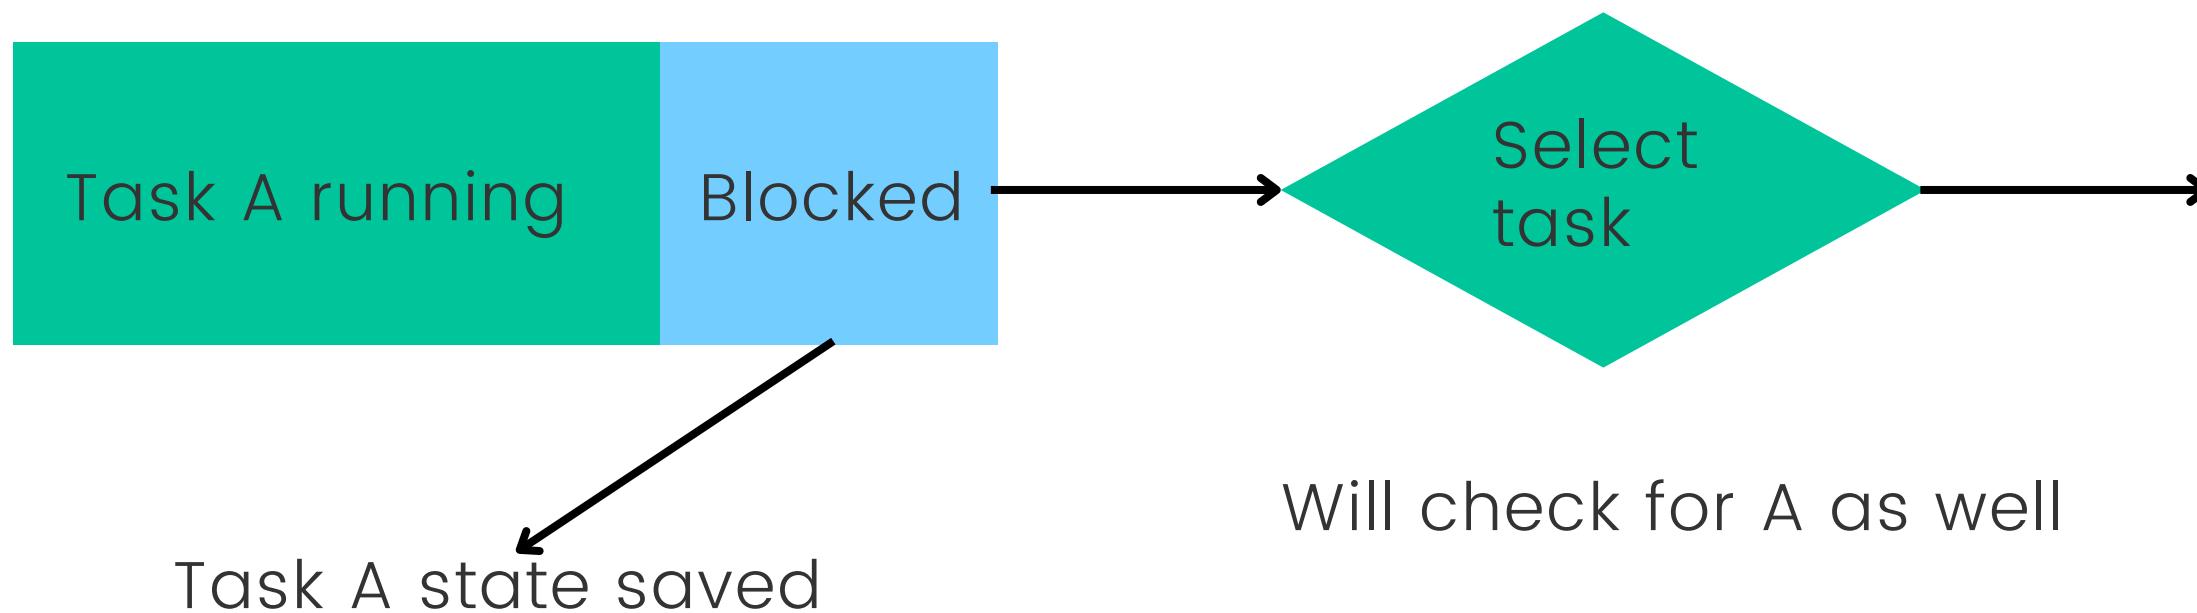
Task A state saved

How to implement concurrency?

Step-3



- Start a task A, and do required computation
- When task A blocks/waits, save state of A
- **Add check in selector to see if A is unblocked, i.e. ready to run again**

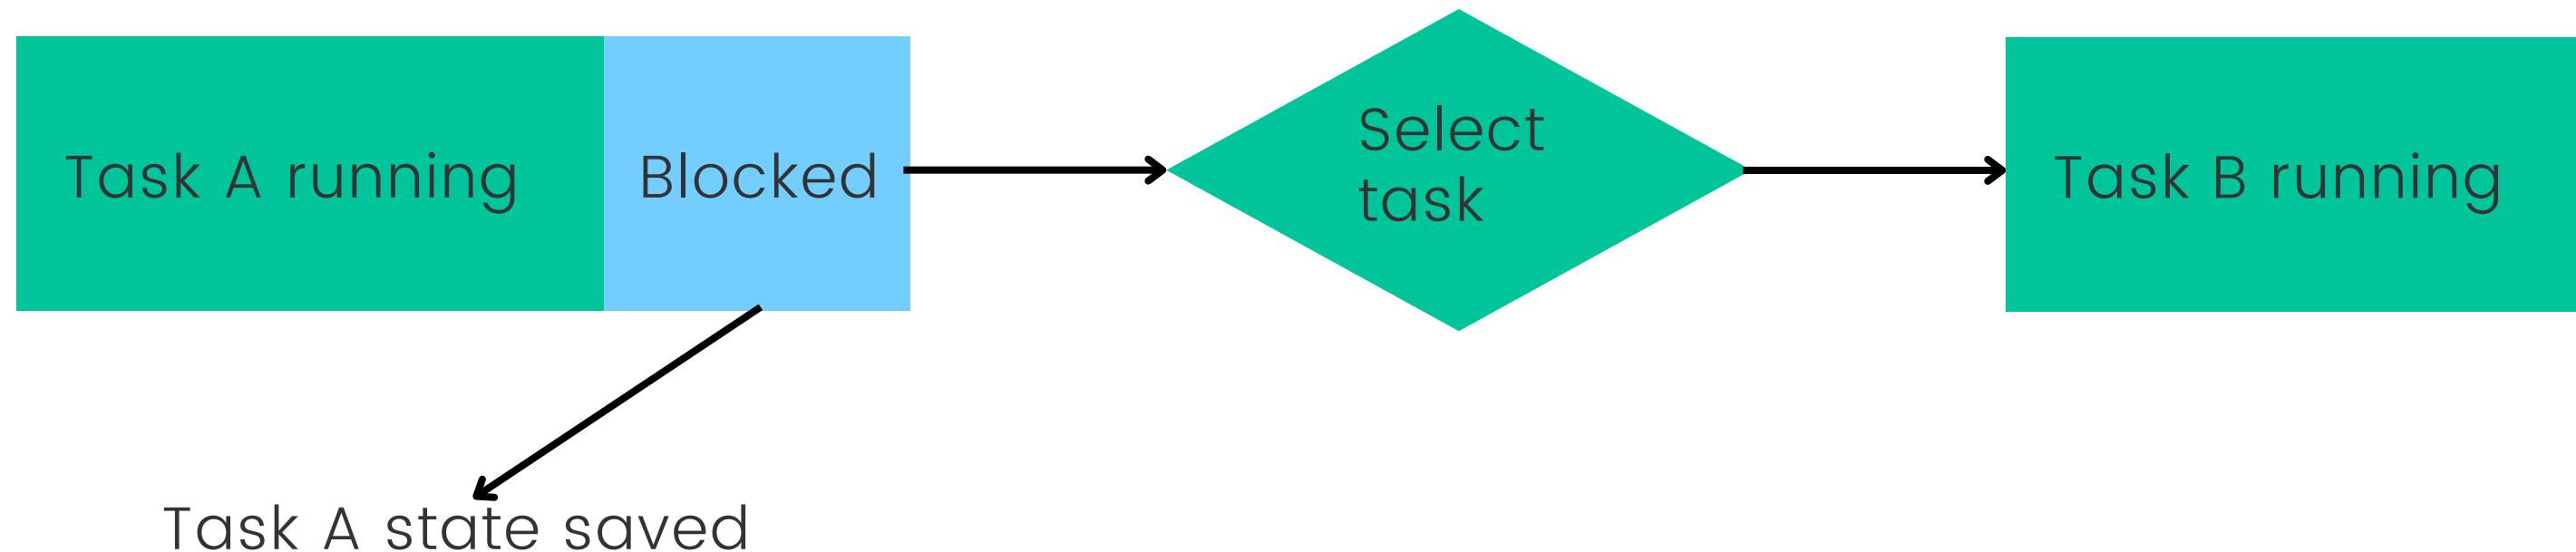


How to implement concurrency?

Step-4 (1/2)



- Start a task A, and do required computation
- When task A blocks/waits, save state of A
- Add check in selector to see if A is unblocked, i.e. ready to run again
- **Check for “ready tasks”, and select one among them to run**

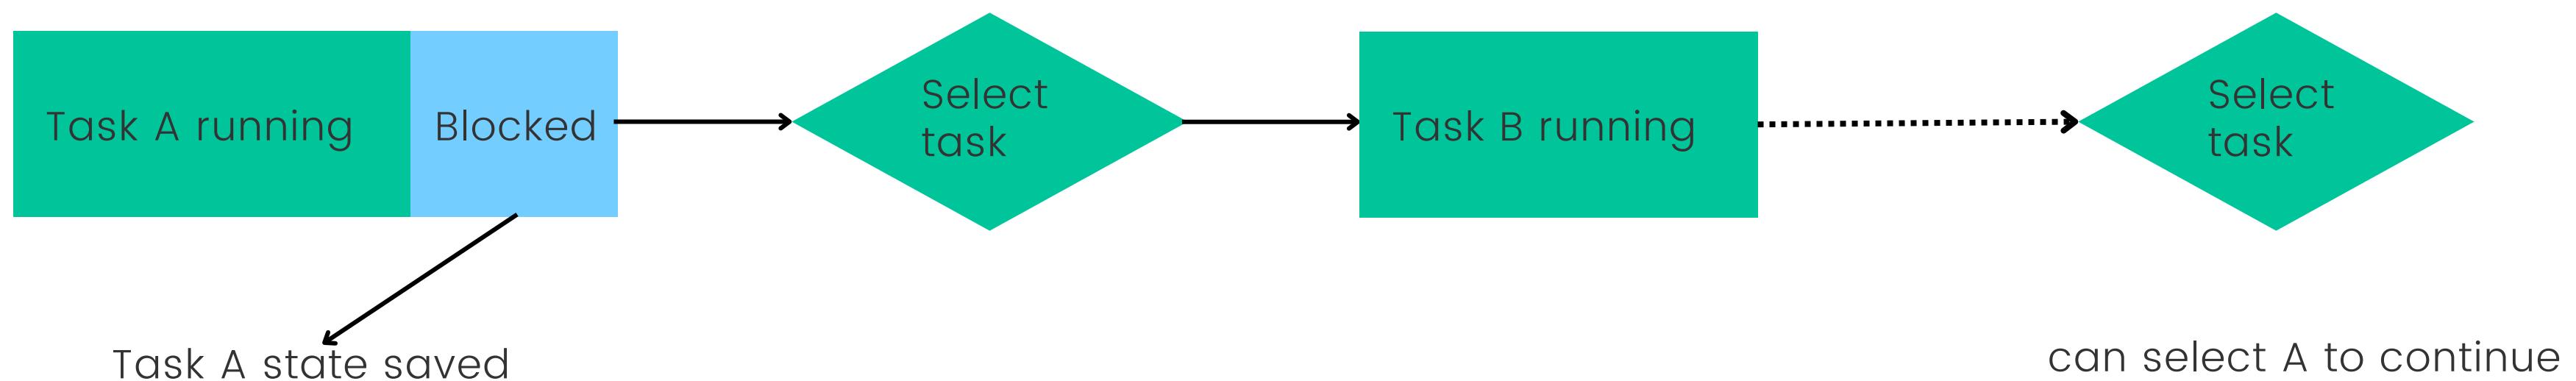


How to implement concurrency?

Step-4 (2/2)



- Start a task A, and do required computation
- When task A blocks/waits, save state of A
- Add check in selector to see if A is unblocked, i.e. ready to run again
- Check for “ready tasks”, and select one among them to run

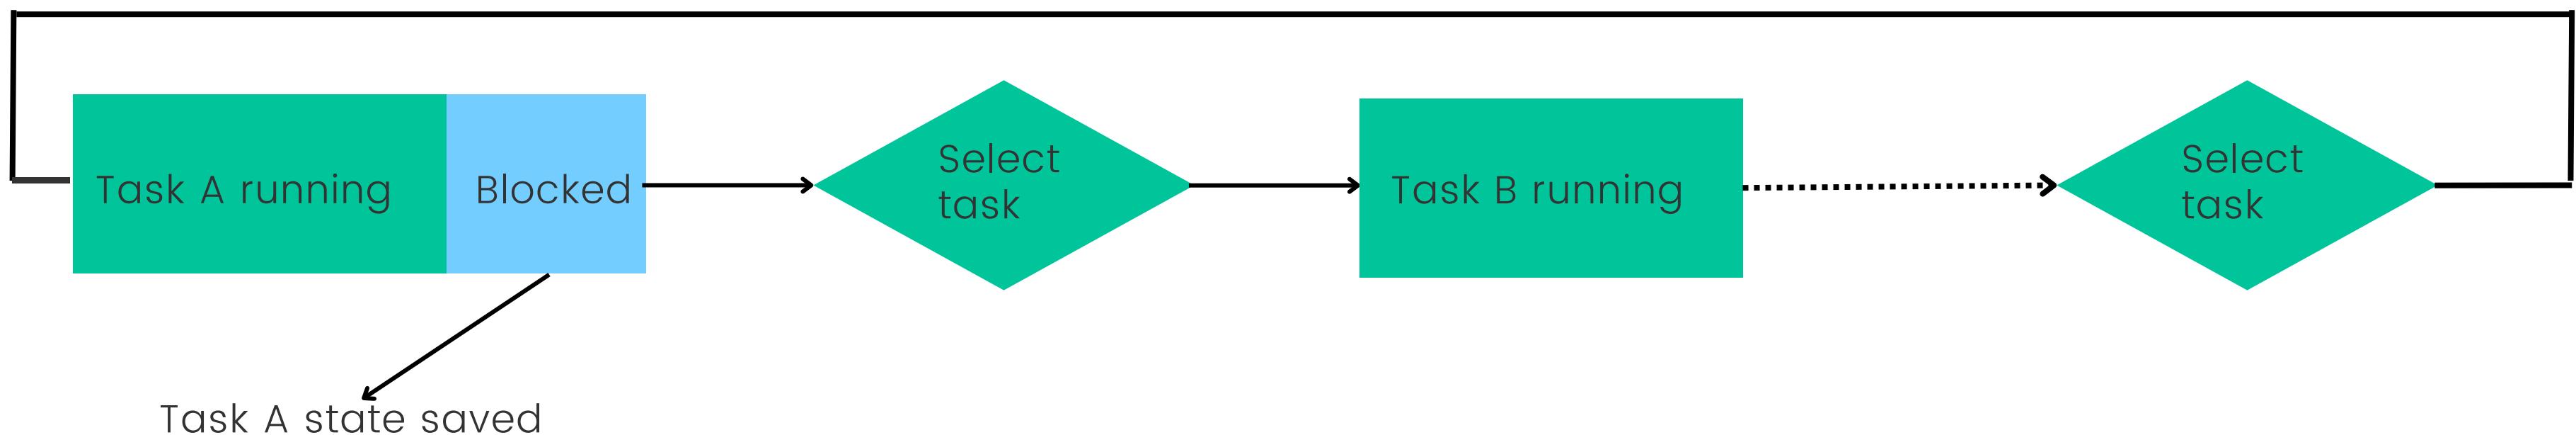


How to implement concurrency?

Step-5



- Start a task A, and do required computation
- When task A blocks/waits, save state of A
- Add check in selector to see if A is unblocked, i.e. ready to run again
- Check for “ready tasks”, and select one among them to run
- **Repeat for each task! (yep this is the event loop)**



So, what's next ?

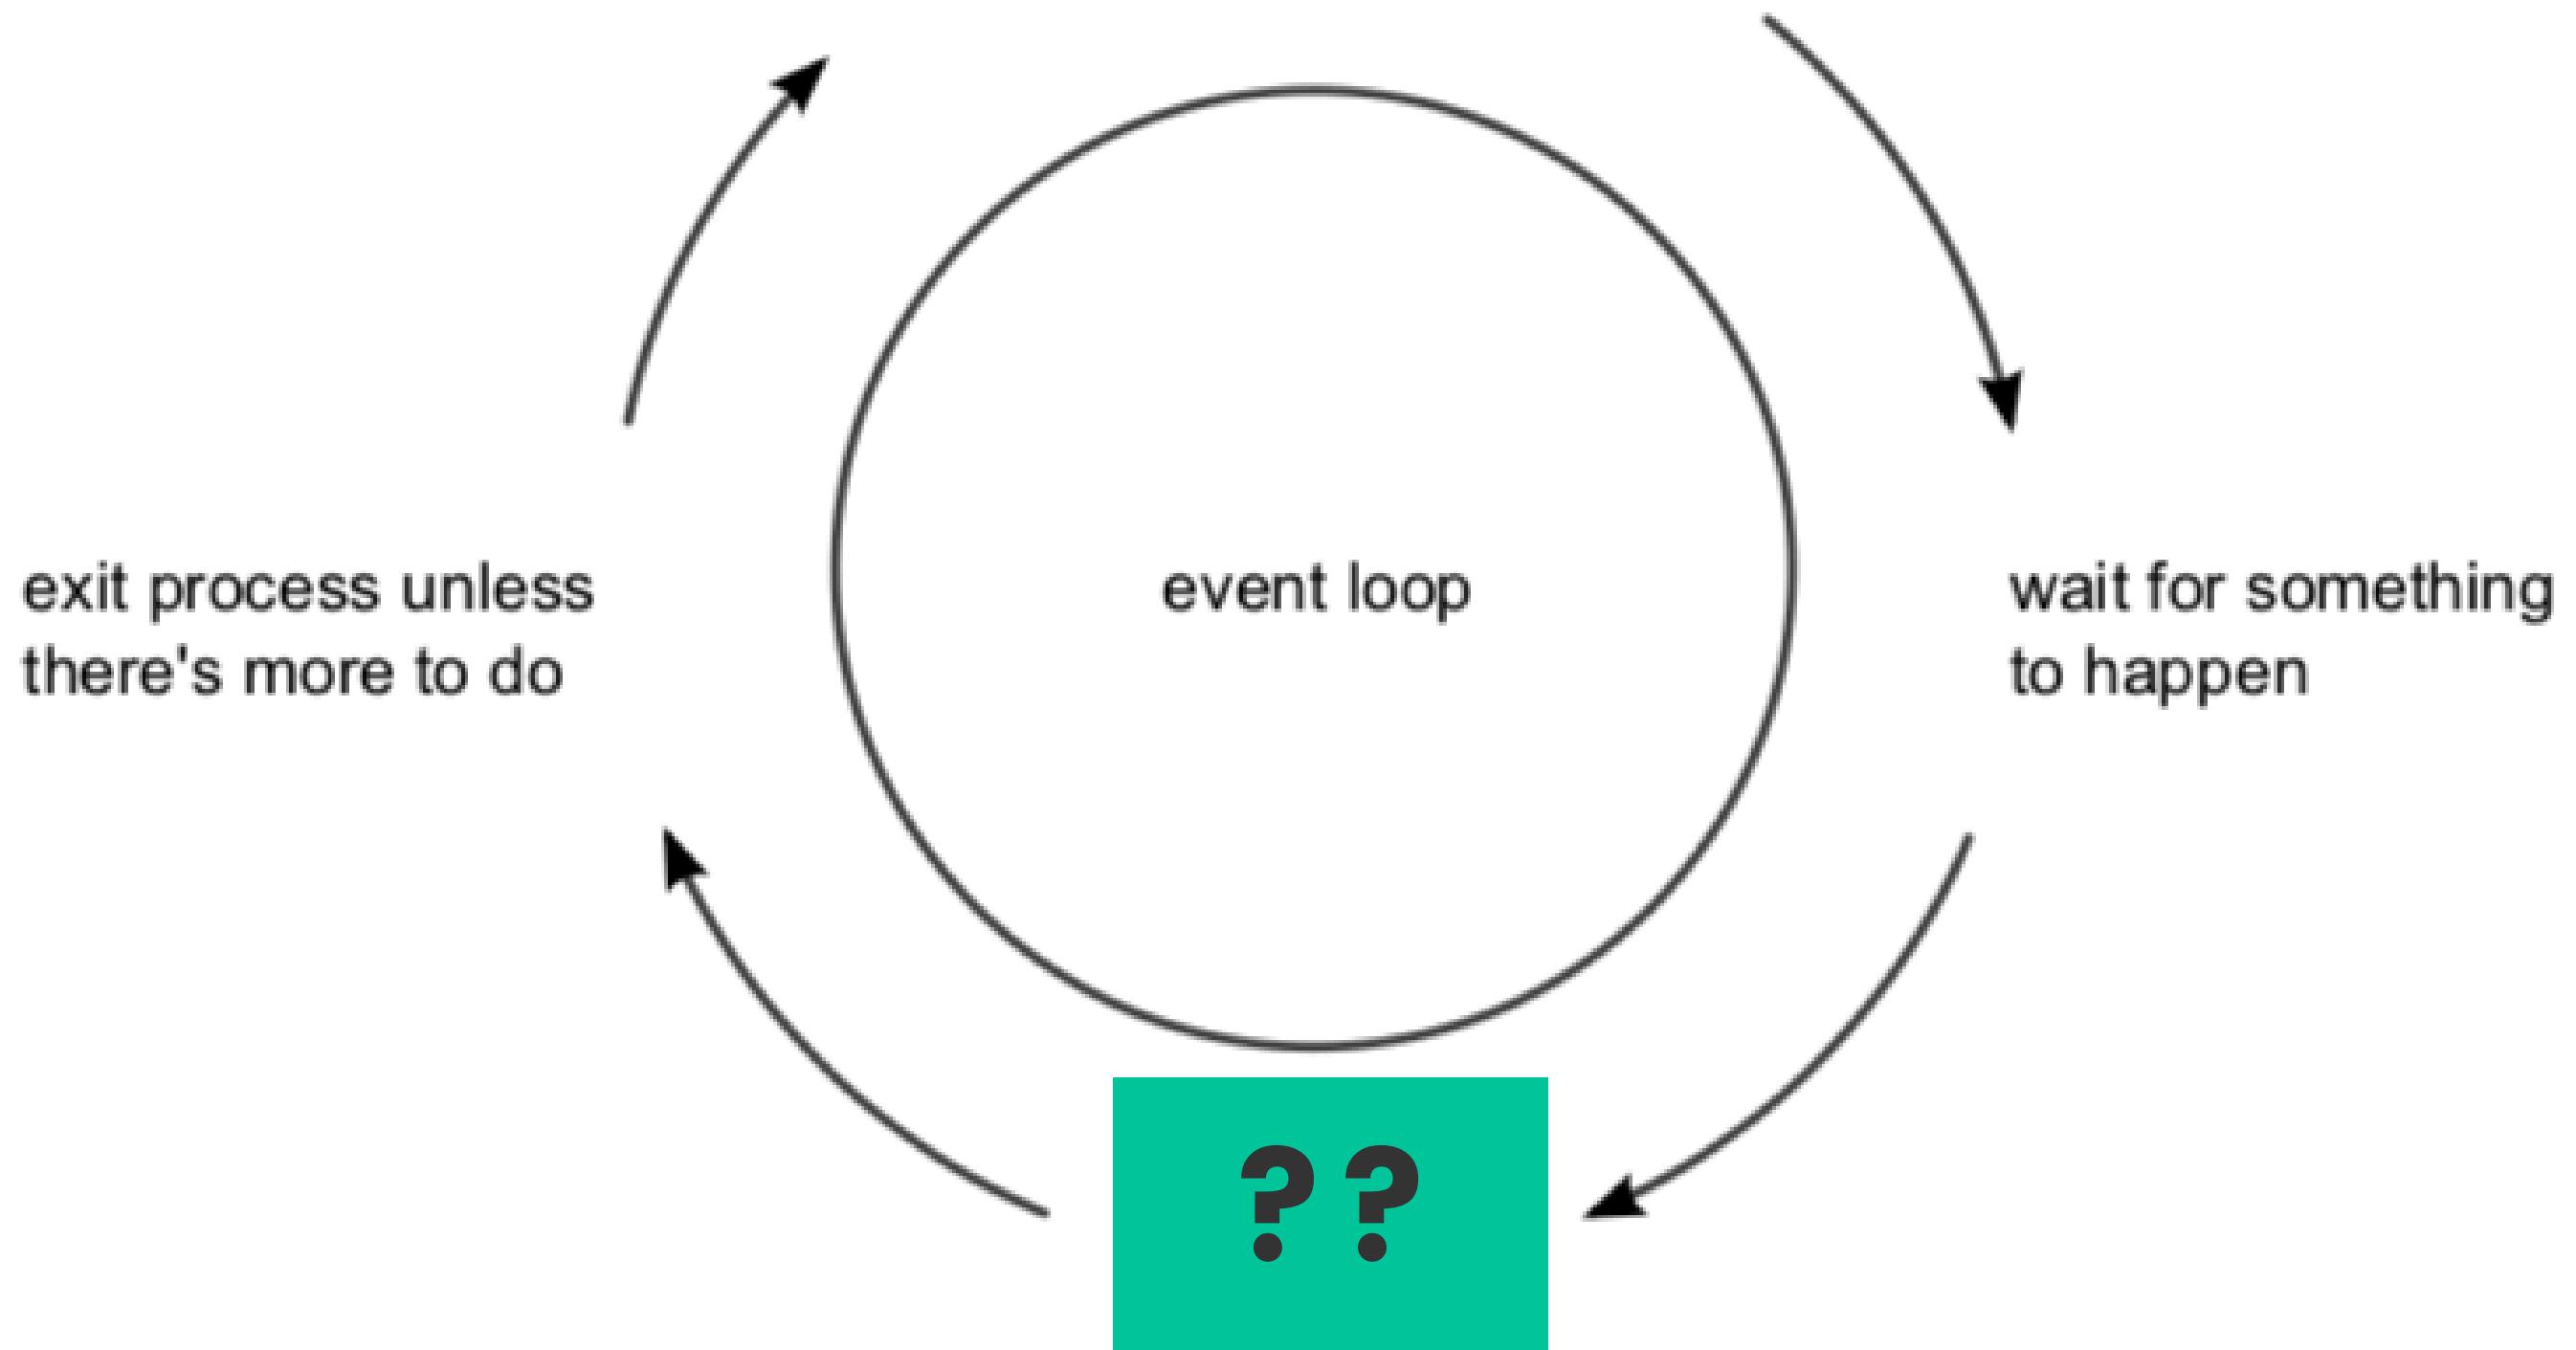


- We have completed “unloading” a function.
- So What's the next Step ?

Hint



Load the program, then → while there's more to do,



So, what's next ?



- We have completed “unloading” a function.
- We have to **reload** the functions ?

~ CALLBACKS

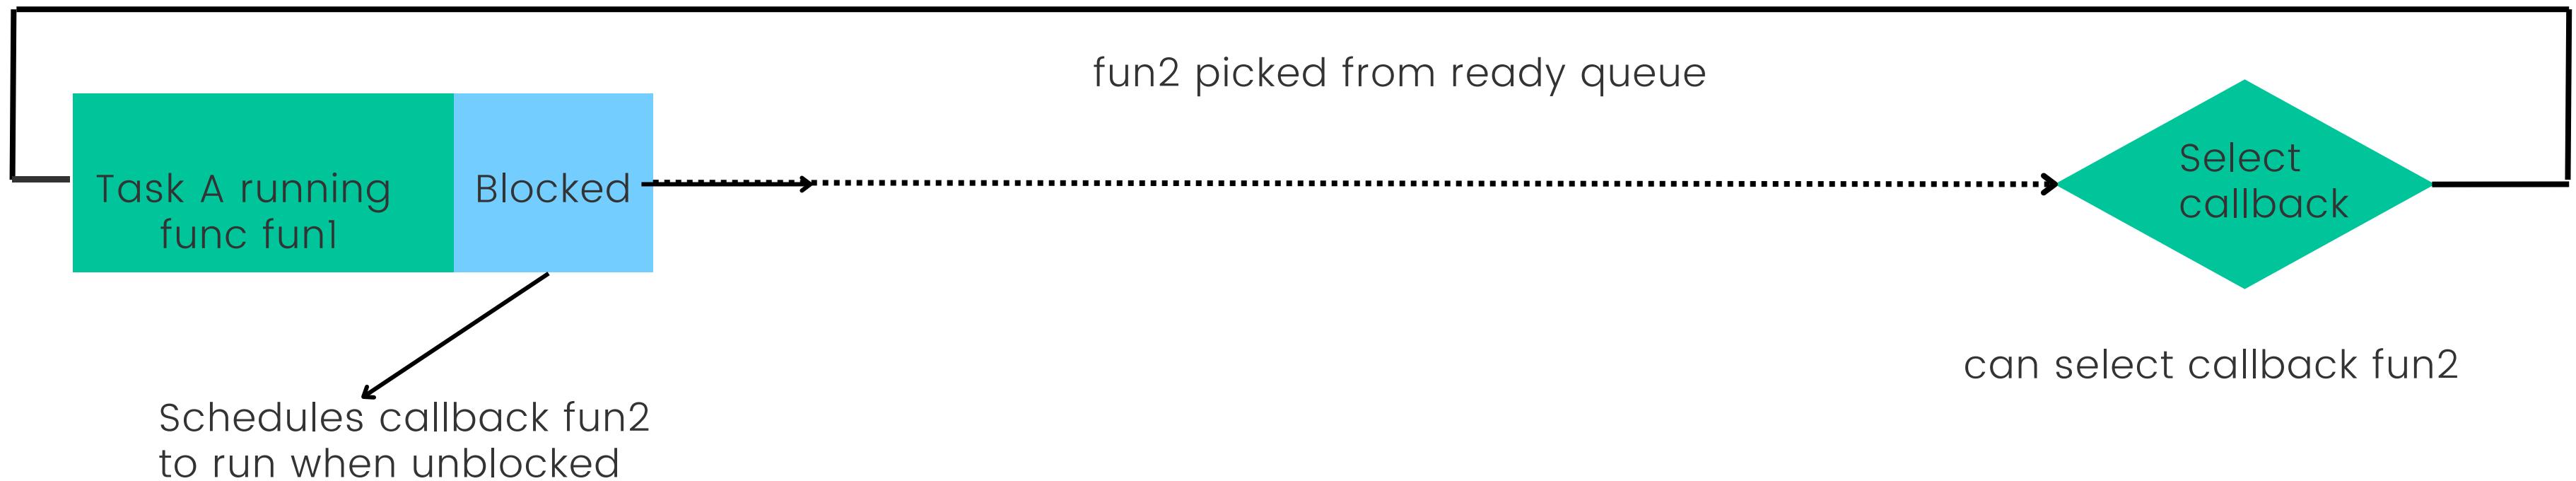
Callbacks



To implement such a design, we use the idea of callbacks. A callback is a piece of code (often a function), used to schedule as a task on the event-loop.

Figure sequence:

func fun1 → schedules callback fun2 → loop calls callback fun2
when unblocked





Event loops & callbacks

- Start a task A, and do required computation
- When task A blocks/waits, save state of A
- Add check in selector to see if A is unblocked, i.e. ready to run again
- Check for “ready tasks”, and select one among them to run
- Repeat for each task! (yep this is the event loop)



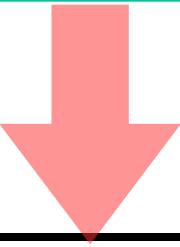
At a high level , we need to define an event loop that does the following while true:

1. Select callback to run from ready queue
2. Run the callback
3. if needed : schedule a new callback



Event loops & callbacks

- Start a task A, and do required computation
- When task A blocks/waits, **save state** of A
- Add **check in selector** to see if A is unblocked, i.e. ready to run again
- Check for “ready tasks”, and select one among them to run
- Repeat for each task! (yep this is the event loop)



At a high level , we need to define an event loop that does the following
while true:

1. Select callback to run from ready queue
2. Run the callback
3. if needed : schedule a new callback

Let's code it ...



EventLoop:

```
ready_tasks = empty queue
blocked_tasks = empty queue

function run():
    while ready_tasks is not empty or \\
        blocked_tasks is not empty:
        check_blocked_tasks()
        if ready_tasks is not empty:
            task = ready_tasks.pop()
            run_task(task)
```



Let's code it ...



```
EventLoop:  
    ready_tasks = empty queue  
    blocked_tasks = empty queue  
  
    function run():  
        while ready_tasks is not empty or \\  
            blocked_tasks is not empty:  
            check_blocked_tasks()  
            if ready_tasks is not empty:  
                task = ready_tasks.pop()  
                run_task(task)
```



```
function check_blocked_tasks():  
    for each task in blocked_tasks:  
        if task is ready:  
            move task from blocked_tasks to ready_tasks  
  
function run_task(task):  
    resume task  
    if task is complete:  
        clean up task  
    else if task is blocked:  
        move task to blocked_tasks  
    else:  
        add task back to ready_tasks
```

Let's code it ...

```
●●●  
EventLoop:  
    ready_tasks = empty queue  
    blocked_tasks = empty queue  
  
    function run():  
        while ready_tasks is not empty or blocked_tasks is not empty:  
            check_blocked_tasks()  
            if ready_tasks is not empty:  
                task = ready_tasks.pop()  
                run_task(task)  
  
    function check_blocked_tasks():  
        for each task in blocked_tasks:  
            if task is ready:  
                move task from blocked_tasks to ready_tasks  
  
    function run_task(task):  
        resume task  
        if task is complete:  
            clean up task  
        else if task is blocked:  
            move task to blocked_tasks  
        else:  
            add task back to ready_tasks
```

Task:

```
function resume():  
    continue execution until next yield point
```

function is_complete():
 return true if task has finished, false otherwise

function is_blocked():
 return true if task is waiting for I/O, false otherwise

function is_ready():
 return true if blocked task can now proceed, false otherwise

Main program:

```
loop = new EventLoop  
add initial tasks to loop.ready_tasks  
loop.run()
```



CURRENT MOOD



Let's start the assignment !!!



CS69201: Computing Lab-1

Systems Assignment 1: Basic and Non-Blocking Execution Environment for C++ Threads

Deadline: October 14, 2024, 11:59 PM

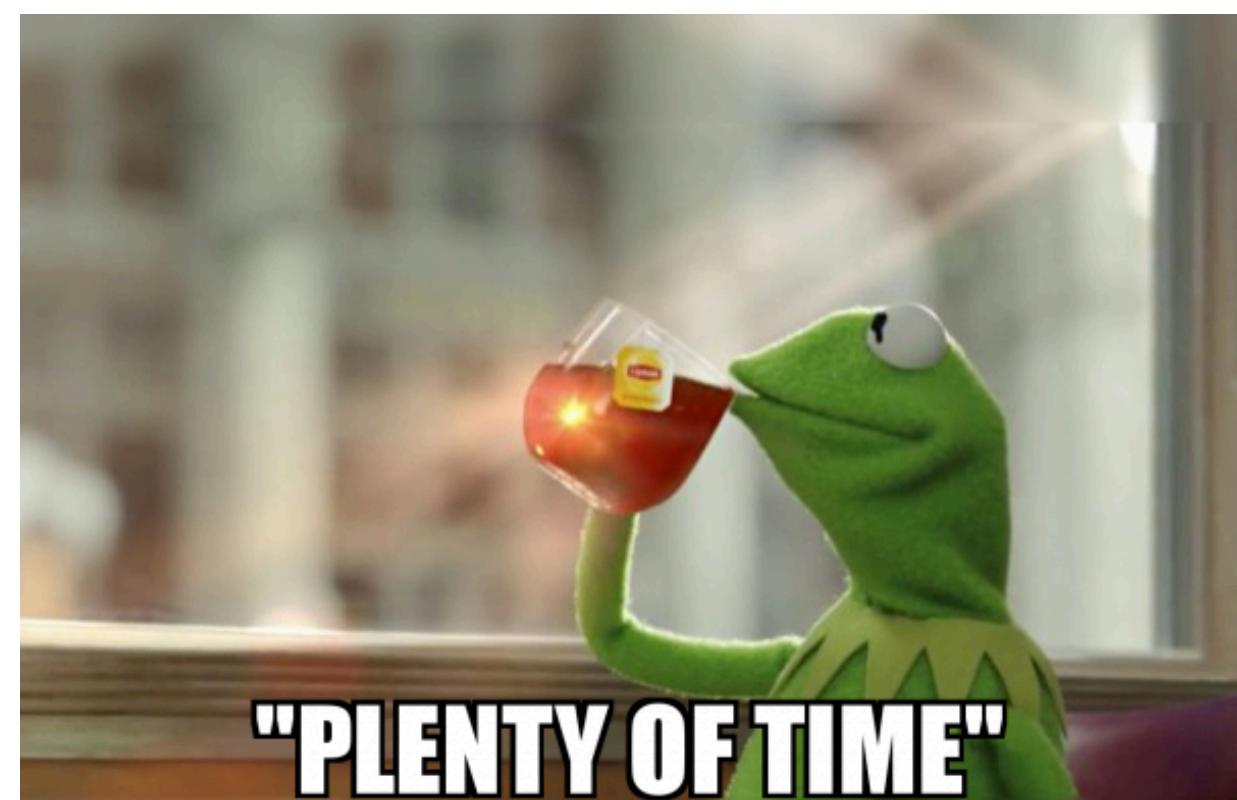
Let's start the assignment !!!



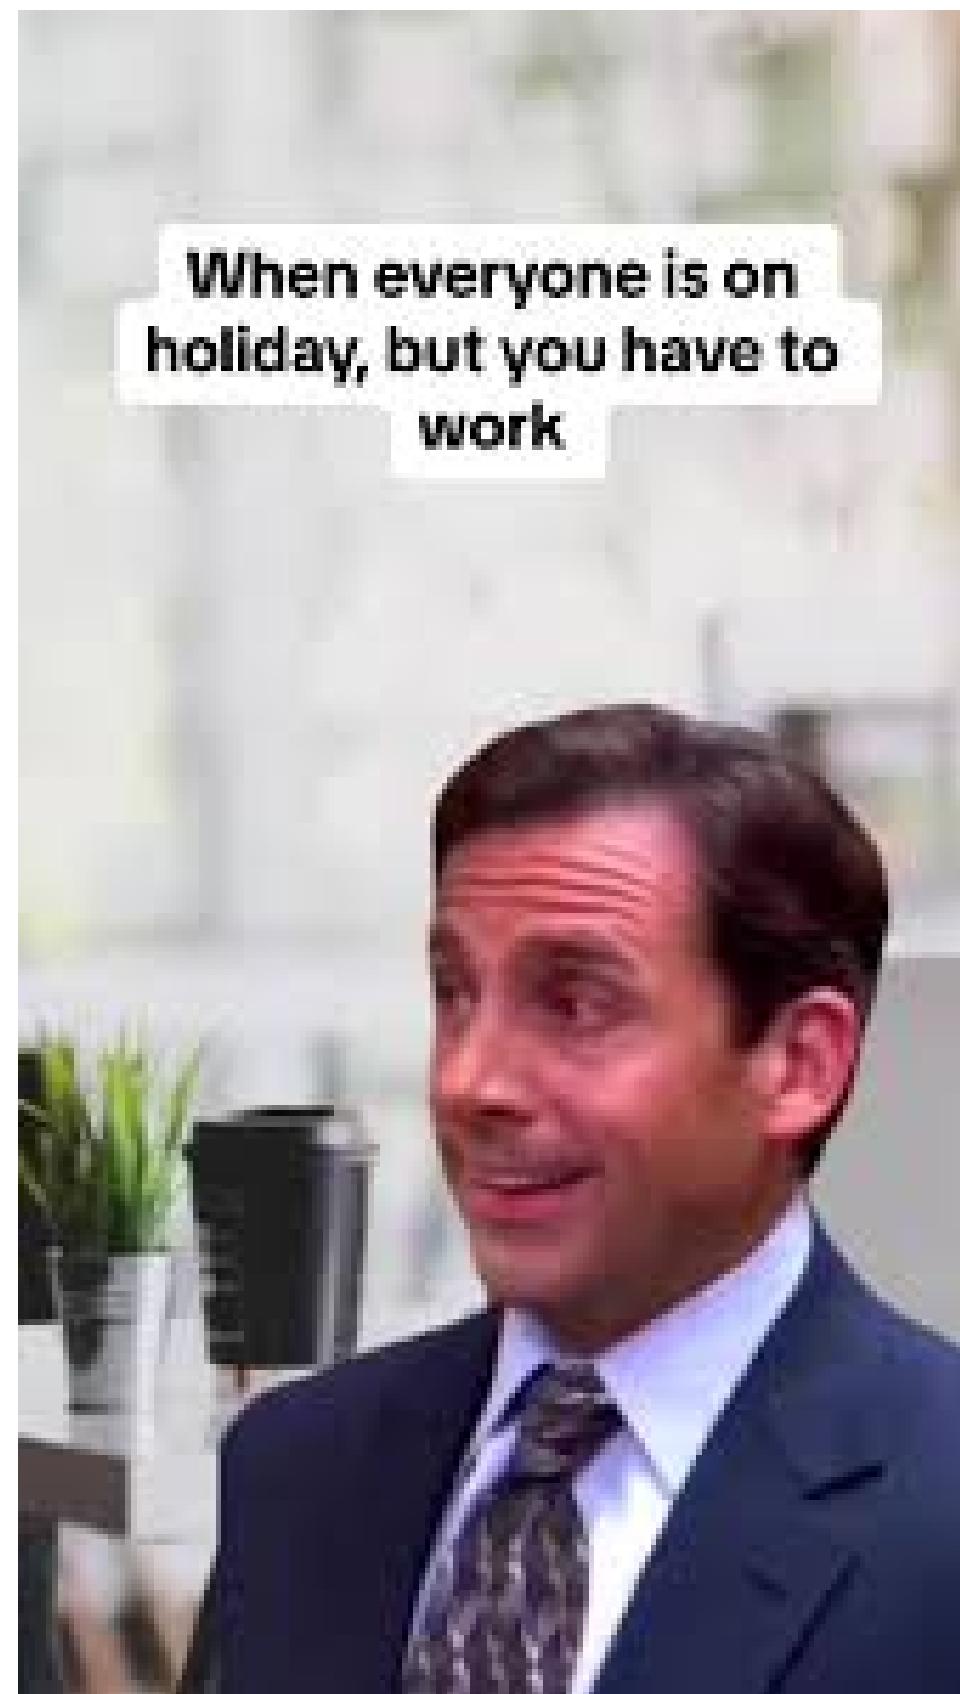
CS69201: Computing Lab-1

Systems Assignment 1: Basic and Non-Blocking Execution Environment for C++ Threads

Deadline: **October 14, 2024, 11:59 PM**







We have another tutorial tomorrow :)



- Date : 2nd October , 2024 [Wednesday]
- Location : CSE 120
- Time : 2-3 PM



**Thank you
for listening!**



Systems Programming: Threading (Part 2)

Computing Lab (cs 69201)
Department of Computer Science and Engineering
IIT Kharagpur

Synchronisation problem



- Do you all know about the critical section problem?
- To get away with it- we use locks.
- Locks :
 - **Mutex Locks**

Mutex Locks



```
// Shared resource
shared_counter = 0

// Mutex lock
mutex = create_mutex()

function increment_counter():
    lock(mutex) // Acquire the lock
        // Critical section
        temp = shared_counter
        temp = temp + 1
        shared_counter = temp

    unlock(mutex) // Release the lock
```



Mutex Locks

```
● ● ●  
  
// Shared resource  
shared_counter = 0  
  
// Mutex lock  
mutex = create_mutex()  
  
function increment_counter():  
    lock(mutex) // Acquire the lock  
        // Critical section  
        temp = shared_counter  
        temp = temp + 1  
        shared_counter = temp  
  
    unlock(mutex) // Release the lock
```

```
● ● ●  
  
// Thread 1  
function thread1_function():  
    for i = 1 to 1000000:  
        increment_counter()  
  
// Thread 2  
function thread2_function():  
    for i = 1 to 1000000:  
        increment_counter()  
  
// Main program  
create_thread(thread1_function)  
create_thread(thread2_function)  
wait_for_threads_to_finish()  
print("Final counter value:", shared_counter)
```

Synchronisation problem



- Do you all know about the critical section problem?
- To get away with it- we use locks.
- Locks :
 - Mutex Locks
 - Semaphores

Semaphores



```
// Initialize a semaphore with 2 permits
semaphore = create_semaphore(2)

function process():
    wait(semaphore) // P operation (down)

    // Critical section
    print("Entering critical section")
    // Simulate some work
    sleep(2 seconds)
    print("Exiting critical section")

    signal(semaphore) // V operation (up)

// Main program
for i = 1 to 5:
    create_thread(process)

wait_for_all_threads_to_finish()
```



Synchronisation problem



- We all know about the critical section problem.
- To ensure atomicity - we use locks
- Locks :
 - Mutex Locks
 - Semaphores
- But what's the **problem** with locks?

Problem with locks



- Performance bottlenecks
- Mutex locks cause context switches
- ... more

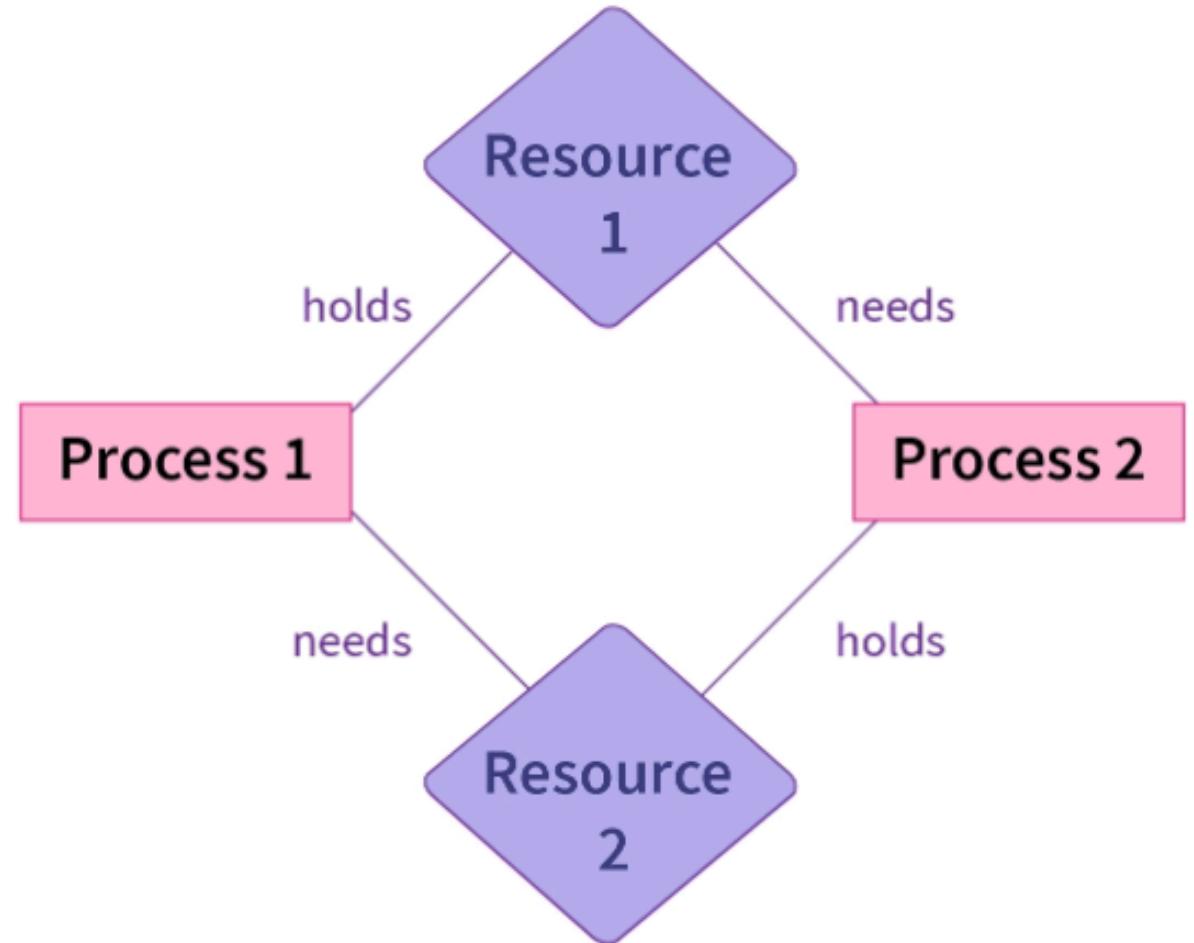


Image credits: scaler.com

Problem with locks



- Performance bottlenecks
- Mutex locks cause context switches
- ... more

So what's the way out ?

Lock-Free Programming



- Thread-safe access to shared data without the use of synchronization primitives such as mutexes.
- Possible but not practical in the absence of hardware support
- Designing them is hard.

Lock-Free Programming



- Thread-safe access to shared data without the use of synchronization primitives such as mutexes.
- Possible but not practical in the absence of hardware support
- Designing them is hard.
- Rather we design **lock-free data structures**. [stack, queue, buffer, map, deque etc.]
- But how ??

Lock-Free Programming



- Thread-safe access to shared data without the use of synchronization primitives such as mutexes.
- Possible but not practical in the absence of hardware support
- Designing them is hard.
- Rather we design **lock-free data structures**. [stack, queue, buffer, map, deque etc.]
- But how ??
 - Well, we will use **lock-free primitives**

Example (1/3)



Structure Node:

```
data  
next (atomic access)
```



Class LockFreeQueue:

```
head // Points to the first node (for dequeue)  
tail // Points to the last node (for enqueue)
```

Constructor:

```
Initialize a dummy node  
Set head and tail to the dummy node
```

Example (2/3)



```
Function Enqueue(value):
    new_node = new Node()
    new_node.data = value
    new_node.next = NULL

Loop:
    current_tail = tail          // Load the current tail pointer
    tail_next = current_tail.next

    // Check if tail pointer is lagging behind (tail.next is not NULL)
    If tail_next != NULL:
        // Move tail forward to catch up
        Compare-And-Swap(tail, current_tail, tail_next)
    Else:
        // Try to append the new node to the end of the queue
        If Compare-And-Swap(current_tail.next, NULL, new_node):
            Break the loop // Successfully linked the new node

    // Now move the tail pointer to the new node
    Compare-And-Swap(tail, current_tail, new_node)
```

Example (3/3)



```
Function Dequeue( ):  
    Loop:  
        current_head = head  
        head_next = current_head.next  
  
        // Check if the queue is empty (head.next is NULL)  
        If head_next == NULL:  
            Return NULL  
  
        // Attempt to move the head forward  
        If Compare-And-Swap(head, current_head, head_next):  
            Return head_next.data    // Successfully dequeued the value
```

Lock-Free APIs in C++



- in C++ we have the atomic library for lock-free programming
- std::atomic provides various atomic data types like:
 - atomic_int
 - atomic_bool
 - atomic_char
 - atomic_intptr_t
 - ...
- provides templatized access to atomic primitives
 - i.e. std::atomic<T>
 - can use with user defined data types (UDT) !!

Lock-Free APIs in c/c++ : Example



```
#include <atomic>
std::atomic<int> counter(0); // Atomic shared counter
void increment() {
    for (int i = 0; i<1000; i++)
        ++counter; // Atomic increment (lock-free)
}

int main() {
    std::vector<std::thread> threads;
    for (int i = 0; i < 100; ++i) {
        threads.push_back(std::thread(increment));
    }
    for (auto& t : threads) {
        t.join();
    }
    std::cout << "Final counter value: " << counter << std::endl;
    return 0;
}
```



Advantages of Lock-Free Program



- No/Less Context-Switches
- Higher CPU frequency and throughput
- No Deadlocks or Priority Inversions
- Faster Multicore Programming

Problems with Lock-Free Program



- Usually, it is hard to write lock-free code
- It is even harder to write correct lock-free code
- Spin-Locks cause heavy memory usage
- Overall performance can go down in some cases compared to mutexes



**Thank you
for listening!**