



Signal Handling, Mutexes, Semaphores and System Calls

**Computing Lab
CS 69201**
**Department of Computer Science and Engineering
IIT Kharagpur**



Preamble

- We delved deep into the basics of System Programming.
- We discussed about Processes, Concurrency, Parallelism, Event Loops, threads, pthreads, IPC, Synchronization, Mutex locks etc.
- We shall ponder upon other aspects - Signal Handling , System calls, Shared Memory.



Signal Handling

- While writing codes, we make some mistakes like accessing unknown memory locations, divide something by 0. This may lead to spurious behaviors during execution. **How to understand?**
- Also, due to some reasons, we would terminate our codes abruptly. That is, we interrupt the system by pressing keys, switching off PC etc. **How can we convey such information?**
- In IPC, we need to communicate with other processes. **How?**

All I Have To Do Is Signal!



Signals

- A signal is a **software generated interrupt** that is sent to a process by the **OS**.
- A signal can report some **exceptional behavior** within the program (such as division by zero), or a signal can report some **asynchronous event** outside the program (such as someone striking an interactive attention key on a keyboard).
- Some examples include :
 - SIGFPE - Floating Point Exception
 - SIGSEGV - The infamous segmentation fault.
 - SIGINT - Interrupt call Ctrl + C
 - SIGILL - Illegal Instruction
 - SIGSTP - Signal Trap Ctrl + Z. (Doing so may lead to process in zombie state - consumes resources for no reason) .

These signals are contained in `<signal.h>`



Signal Handler

- A signal handler is a function which is called by the target environment when the corresponding signal occurs.
- **signal()** is a function which handles signals.
- The signal() call takes two parameters: the signal in question, and an action to take when that signal is raised.



Signal Handler

```
1  /* A C program that does not terminate when Ctrl+C is pressed */
2  #include <stdio.h>
3  #include <signal.h>
4
5  /* Signal Handler for SIGINT */
6  void sigintHandler(int sig_num)
7  {
8      /* Reset handler to catch SIGINT next time.
9       Refer http://en.cppreference.com/w/c/program/signal */
10     signal(SIGINT, sigintHandler);
11     printf("\n Cannot be terminated using Ctrl+C \n");
12     fflush(stdout);
13 }
14
15 int main ()
16 {
17     /* Set the SIGINT (Ctrl-C) signal handler to sigintHandler
18     Refer http://en.cppreference.com/w/c/program/signal */
19     signal(SIGINT, sigintHandler);
20
21     /* Infinite loop */
22     while(1)
23     {
24     }
25     return 0;
26 }
```

Output: When Ctrl+C was pressed two times

Cannot be terminated using Ctrl+C

Cannot be terminated using Ctrl+C



pThreads

- For UNIX systems, implementations of threads that adhere to the IEEE POSIX 1003.1c standard are Pthreads.
- The primary motivation behind Pthreads is improving program performance
- Can be created with much less OS overhead & Needs fewer system resources to run.



pThreads Library

- Programs must include the file pthread.h
- Programs must be linked with the pthread library
 - gcc -pthread main.c -o main
- Some functions :
 - `int pthread_create(pthread_t * thread, const pthread_attr_t * attr, void * (*start_routine)(void *), void *arg);`: Creates a new thread
 - thread: pointer to an unsigned integer value that returns the thread id of the thread created.
 - attr: pointer to a structure that is used to define thread attributes like detached state, scheduling policy, stack address, etc. Set to NULL for default thread attributes.
 - start_routine: pointer to a subroutine that is executed by the thread. The return type and parameter type of the subroutine must be of type void *. The function has a single attribute but if multiple values need to be passed to the function, a struct must be used.
 - arg: pointer to void that contains the arguments to the function defined in the earlier argument



pThreads Library

➤ Some functions :

- void pthread_exit(void *retval): Terminates the calling thread
 - This method accepts a mandatory parameter retval which is the pointer to an integer that stores the return status of the thread terminated.
 - The scope of this variable must be global so that any thread waiting to join this thread may read the return status.
- pthread_join(): Causes the calling thread to wait for another thread to terminate



pThreads Library



```
#include <pthread.h>

void *thread_function(void *arg) {
    // Thread code here
    return NULL;
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, thread_function, NULL);
    pthread_join(thread, NULL);
    return 0;
}
```



Mutexes

- A mutex (mutual exclusion) is a synchronization object that is used to protect shared data from concurrent access by multiple threads.
- A mutex is a binary semaphore that can be in one of two states: locked or unlocked.
- Only one thread can hold a mutex at a time.



Mutexes

- To acquire a mutex, you call the `pthread_mutex_lock()` function.
 - `int pthread_mutex_init(pthread_mutex_t *restrict mutex, const pthread_mutexattr_t *restrict attr)` : Creates a mutex, referenced by mutex, with attributes specified by attr.
 - Returns 0 if successful else -1
- To release a mutex, you call the `pthread_mutex_unlock()` function.
 - `int pthread_mutex_lock(pthread_mutex_t *mutex)` : Locks a mutex object, which identifies a mutex. If the mutex is already locked by another thread, the thread waits for the mutex to become available.
 - Returns 0 if successful else -1.
- If a thread tries to acquire a mutex that is already locked, it will block until the mutex is released.



Mutex

```
1 #include <pthread.h>
2 #include <stdio.h>
3
4 int shared_data = 0;
5 pthread_mutex_t mutex;
6
7 void *thread_func(void *arg) {
8     for (int i = 0; i < 1000; i++) {
9         pthread_mutex_lock(&mutex);
10        shared_data++;
11        pthread_mutex_unlock(&mutex);
12    }
13    return NULL;
14 }
```

```
16 int main() {
17     pthread_t thread1, thread2;
18     pthread_mutex_init(&mutex, NULL);
19
20     pthread_create(&thread1, NULL, thread_func, NULL);
21     pthread_create(&thread2, NULL, thread_func, NULL);
22
23     pthread_join(thread1, NULL); []
24     pthread_join(thread2, NULL);
25
26     printf("Shared []"
27           "data: %d\n", shared_data);
28
29     return 0;
30 }
```



Semaphores

- A synchronization mechanism used to control access to shared resources.
- A non-negative integer.
- Operations:
 - **wait(sem)**: Decrements the semaphore value. If the value becomes negative, the process is blocked.
 - **signal(sem)**: Increments the semaphore value. If there are any blocked processes, one of them is unblocked.
- Types of Semaphores
 - **Counting Semaphore**: Can have any non-negative integer value. Used to control access to a resource with a limited number of available instances.
 - **Binary Semaphore**: Can only have the values 0 and 1. Used to protect critical sections of code.



Semaphores

```
1 #include <semaphore.h>
2
3 sem_t resource_semaphore;
4
5 void* producer(void* arg) {
6     while (1) {
7         // Produce an item
8         sem_wait(&resource_semaphore);
9         // Place the item in the buffer
10        sem_post(&resource_semaphore);
11    }
12 }
13
14 void* consumer(void* arg) {
15     while (1) {
16         sem_wait(&resource_semaphore);
17         // Take an item from the buffer
18         sem_post(&resource_semaphore);
19         // Consume the item
20    }
21 }
```



Semaphores

- Operations (include “#include<semaphore.h>”):
 - void sem_init(sem_t *sem, int pshared, unsigned int value):
 - sem : Specifies the semaphore to be initialized.
 - pshared : This argument specifies whether or not the newly initialized semaphore is shared between processes or between threads. A non-zero value means the semaphore is shared between processes and a value of zero means it is shared between threads.
 - value : Specifies the value to assign to the newly initialized semaphore
 - int sem_wait(sem_t *sem): Decrements the semaphore value. If the value becomes negative, the process is blocked. sem : Specifies the semaphore to be initialized.
 - int sem_post(sem_t *sem): Increments the semaphore value. If there are any blocked processes, one of them is unblocked.
 - sem_destroy(sem_t *mutex); Destroys semaphore
- Types of Semaphores
 - Counting Semaphore: Can have any non-negative integer value. Used to control access to a resource with a limited number of available instances.
 - Binary Semaphore: Can only have the values 0 and 1. Used to protect critical sections of code.



Binary Semaphores

```
1 #include <stdio.h>
2 #include <semaphore.h>
3
4 sem_t mutex;
5
6 void *thread_func(void *arg) {
7     int i;
8
9     for (i = 0; i < 5; i++) {
10         sem_wait(&mutex);
11         printf("Thread %d: accessing critical section\n", (int)arg);
12         // Critical section code here
13         sem_post(&mutex);
14     }
15
16     return NULL;
17 }
18 }
```

```
1 v int main() {
2     pthread_t thread1, thread2;
3
4     // Initialize the semaphore
5     sem_init(&mutex, 0, 1);
6
7     // Create threads
8     pthread_create(&thread1, NULL, thread_func, (void*)1);
9     pthread_create(&thread2, NULL, thread_func, (void*)2);
10
11    // Wait for threads to finish
12    pthread_join(thread1, NULL);
13    pthread_join(thread2, NULL);
14
15    // Destroy the semaphore
16    sem_destroy(&mutex);
17
18    return 0;
19 }
```



Counting Semaphores

```
1 #include <stdio.h>
2 #include <semaphore.h>
3
4 sem_t resource_count;
5
6 void *producer_func(void *arg) {
7     int i;
8
9     for (i = 0; i < 10; i++) {
10         sem_wait(&resource_count);
11         // Produce a resource
12         printf("Producer %d: produced a resource\n", (int)arg);
13     }
14
15     return NULL;
16 }
```

```
18 void *consumer_func(void *arg) {
19     int i;
20
21     for (i = 0; i < 10; i++) {
22         sem_wait(&resource_count);
23         // Consume a resource
24         printf("Consumer %d: consumed a resource\n", (int)arg);
25         sem_post(&resource_count);
26     }
27
28     return NULL;
29 }
30 }
```



Counting Semaphores

```
1 int main() {
2     pthread_t producers[3], consumers[2];
3     int i;
4
5     // Initialize the semaphore with an initial count of 5
6     sem_init(&resource_count, 0, 5);
7
8     // Create producer threads
9     for (i = 0; i < 3; i++) {
10         pthread_create(&producers[i], NULL, producer_func, (void*)i);
11     }
12
13     // Create consumer threads
14     for (i = 0; i < 2; i++) {
15         pthread_create(&consumers[i], NULL, consumer_func, (void*)i);
16     }
17 }
```

```
18     // Wait for threads to finish
19     for (i = 0; i < 3; i++) {
20         pthread_join(producers[i], NULL);
21     }
22     for (i = 0; i < 2; i++) {
23         pthread_join(consumers[i], NULL);
24     }
25
26     // Destroy the semaphore
27     sem_destroy(&resource_count);
28
29     return 0;
30 }
```



Shared Memory

- A technique where multiple processes can access the same memory region. This allows for efficient communication and data sharing between processes.
- Operations (include the line “#include<sys/shm.h>”):
 - `shmget(IPC_PRIVATE, sizeof(int) * BUFFER_SIZE, IPC_CREAT | 0666)`:
 - Creates a shared memory segment with a unique identifier (IPC_PRIVATE).
 - Allocates memory for the buffer with a size of `sizeof(int) * BUFFER_SIZE`.
 - Sets the permissions for the shared memory segment to 0666 (read and write permissions for owner, group, and others).
 - `buffer = shmat(shmid, NULL, 0)`
 - Attaches the shared memory segment to the process's address space.
 - Returns a pointer to the attached memory, which is stored in the `buffer` variable.
 - `shmdt(buffer)`: Detach the shared memory segment from the process's address space.
 - `shmctl(shmid, IPC_RMID, NULL)`: Remove the shared memory segment from the system.



Shared Memory - Producer Consumer

```
1 #include <pthread.h>
2 #include <semaphore.h>
3 #include <sys/shm.h>
4
5 #define BUFFER_SIZE 10
6
7 int *buffer;
8 int in = 0, out = 0;
9
10 sem_t empty_slots;
11 sem_t full_slots;
12 pthread_mutex_t mutex;
13
14 void* producer(void* arg) {
15     while (1) {
16         sem_wait(&empty_slots);
17         pthread_mutex_lock(&mutex);
18         buffer[in] = /* produce an item */;
19         in = (in + 1) % BUFFER_SIZE;
20         pthread_mutex_unlock(&mutex);
21         sem_post(&full_slots);
22     }
23 }
```

```
25 void* consumer(void* arg) {
26     while (1) {
27         sem_wait(&full_slots);
28         pthread_mutex_lock(&mutex);
29         int item = buffer[out];
30         out = (out + 1) % BUFFER_SIZE;
31         pthread_mutex_unlock(&mutex);
32         sem_post(&empty_slots);
33         // consume the item
34     }
35 }
36 int main() {
37     // Create shared memory segment
38     int shmid = shmget(IPC_PRIVATE, sizeof(int) * BUFFER_SIZE, IPC_CREAT | 0666);
39     buffer = shmat(shmid, NULL, 0);
40
41     // Initialize semaphores and mutex
42     sem_init(&empty_slots, 0, BUFFER_SIZE);
43     sem_init(&full_slots, 0, 0);
44     pthread_mutex_init(&mutex, NULL);
45
46     // Create producer and consumer threads
47     pthread_t producer_thread, consumer_thread;
```



Shared Memory - Producer Consumer

```
38     int shmid = shmget(IPC_PRIVATE, sizeof(int) * BUFFER_SIZE, IPC_CREAT | 0666);
39     buffer = shmat(shmid, NULL, 0);
40
41     // Initialize semaphores and mutex
42     sem_init(&empty_slots, 0, BUFFER_SIZE);
43     sem_init(&full_slots, 0, 0);
44     pthread_mutex_init(&mutex, NULL);
45
46     // Create producer and consumer threads
47     pthread_t producer_thread, consumer_thread;
48     pthread_create(&producer_thread, NULL, producer, NULL);
49     pthread_create(&consumer_thread, NULL, consumer, NULL);
50
51
52     // Join threads and detach shared memory
53     pthread_join(producer_thread, NULL);
54     pthread_join(consumer_thread, NULL);
55     shmdt(buffer);
56     shmctl(shmid, IPC_RMID, NULL);
57
58     return 0;
59 }
```



System Calls

- A system call is a programmatic way in which a computer program requests a service from the kernel of the operating system it is executed on.
- A system call is a way for programs to interact with the operating system.
- A computer program makes a system call when it requests the operating system's kernel.
- Features of System Calls :
 - Provide a well-defined interface between user programs and the operating system.
 - Access privileged operations that are not available to normal user programs.
 - Switch to Kernel Mode
 - Context Switch
 - Error Handling
 - Synchronization



System Calls

- Procedure of System Call Invocation :
 - Users need special resources
 - The program makes a system call request
 - Operating system sees the system call
 - The operating system performs the operations
 - Operating system give control back to the program
- Examples : fork(), pipe()



fork()

- Creates a new process as a child process of the calling process (parent)
- Both have similar code segments
- The child gets a copy of the parents data segment at the time of forking
- Returns the following :
 - Negative Value: The creation of a child process was unsuccessful.
 - Zero: Returned to the newly created child process.
 - Positive value: Returned to parent or caller. The value contains the process ID of the newly created child process.



fork()

```
1 #include <stdio.h>
2 #include <sys/types.h>
3 #include <unistd.h>
4 int main()
5 {
6
7     // make two process which run same
8     // program after this instruction
9     pid_t p = fork();
10    if(p<0){
11        perror("fork fail");
12        exit(1);
13    }
14    printf("Hello world!, process_id(pid) = %d \n",getpid());
15    return 0;
16 }
```

Output

```
Hello world!, process_id(pid) = 31
Hello world!, process_id(pid) = 32
```



fork()

```
1 #include <stdio.h>
2 #include <sys/types.h>
3 #include <unistd.h>
4 int main()
5 {
6     fork();
7     fork();
8     fork();
9     printf("hello\n");
10    return 0;
11 }
```

Output

```
hello
hello
hello
hello
hello
hello
hello
hello
```



pipe()

- The pipe() system call creates a pipe that can be shared between processes
- It returns two file descriptors, one for reading from the pipe, the other for writing into the pipe.
- File descriptors (fd) is an integer that uniquely identifies an open file of the process.
- 0 for stdin, 1 for stdout, 2 for stderr.



pipe()

```
1 #include <stdio.h>
2 #include <unistd.h> /* Include this file to use pipes */
3 #define BUFSIZE 80
4 main()
5 {
6     int fd[2], n=0, i;
7     char line[BUFSIZE];
8     pipe(fd); /* fd[0] is for reading, fd[1] is for writing */ if (fork() == 0) {
9         close(fd[0]); /* The child will not read */
10
11    for (i=0; i < 10; i++) {
12        sprintf(line,"%d",n);
13        write(fd[1], line, BUFSIZE);
14        printf("Child writes: %d\n",n); n++; sleep(2);
15    }
16    else {
17        close(fd[1]); /* The parent will not write */
18        for (i=0; i < 10; i++) {
19            read(fd[0], line, BUFSIZE);
20            sscanf(line,"%d",&n);
21            printf("\t\t\t Parent reads: %d\n",n);
22        }
    }
```

Thank You